


Latent Variable Models

Overview of Latent Variable Models

- **Motivation:** Find “latent” (unobserved) structure in neural population activity
- Latents can be discrete or continuous

Overview of Latent Variable Models


- **Motivation:** Find “latent” (unobserved) structure in neural population activity
- Latents can be discrete or **continuous**



Cunningham & Yu, 2014



Overview of Latent Variable Models

- **Motivation:** Find “latent” (unobserved) structure in neural population activity
- Latents can be discrete or **continuous**




Overview of Latent Variable Models

- **Motivation:** Find “latent” (unobserved) structure in neural population activity
- Latents can be discrete or **continuous**




Overview of Latent Variable Models

- **Motivation:** Find “latent” (unobserved) structure in neural population activity
- Latents can be discrete or **continuous**




Overview of Latent Variable Models

- **Motivation:** Find “latent” (unobserved) structure in neural population activity
- Latents can be discrete or **continuous**



Overview of Latent Variable Models


- **Motivation:** Find “latent” (unobserved) structure in neural population activity
- Latents can be discrete or **continuous**



Cunningham & Yu, 2014

Overview of Latent Variable Models


- **Motivation:** Find “latent” (unobserved) structure in neural population activity
- Latents can be **discrete** or continuous




Overview of Workshop

- **Day 1:**
 - Intro to LVMs
 - Factor Analysis (FA)
 - Gaussian Processes (GPs)
 - Gaussian Process Factor Analysis (GPFA)
 - Hidden Markov Models (HMMs)
- **Day 2:**
 - Linear Dynamical Systems (LDSs)
 - Variational Autoencoders (VAEs)
 - Many variants
- **Day 3:**
 - Recap (mini-presentations?) and final questions

Intro to Latent Variable Models




Intro to Latent Variable Models




- Two parts of a LVM
 - Prior: $z \sim p(z)$
- Conditional probability of observed data: $x|z \sim p(x|z)$

Intro to Latent Variable Models




Intro to Latent Variable Models

- Probability of observed data, $p(x)$, is:




Intro to Latent Variable Models



- Probability of observed data, $p(x)$, is:
- For discrete latents:

$$p(x) = \sum_{i=1}^m p(x|z = z_i)p(z = z_i)$$

Intro to Latent Variable Models




- Probability of observed data, $p(x)$, is:
 - For discrete latents:

$$p(x) = \sum_{i=1}^m p(x|z = z_i)p(z = z_i)$$

- For continuous latents:


$$p(x) = \int p(x|z)p(z)dz$$

Intro to Latent Variable Models: Goals



Intro to Latent Variable Models: Goals


- Recognition/Inference



$$p(z|x)$$

Intro to Latent Variable Models: Goals


- **Recognition/Inference**



$$p(z|x) = \frac{p(x|z)p(z)}{p(x)}$$

Intro to Latent Variable Models: Goals

- **Recognition/Inference**



$$p(z|x) = \frac{p(x|z)p(z)}{p(x)}$$


- **Model Fitting**

- Model including parameters is actually:

$$p(x, z|\theta) = p(x|z, \theta)p(z|\theta)$$

Intro to Latent Variable Models: Goals

- **Recognition/Inference**



$$p(z|x) = \frac{p(x|z)p(z)}{p(x)},$$

- **Model Fitting**

- Model including parameters is actually:


$$p(x, z|\theta) = p(x|z, \theta)p(z|\theta)$$

- Learning parameters by maximum likelihood:


$$\hat{\theta} = \arg \max_{\theta} p(x|\theta) = \arg \max_{\theta} \int p(x, z|\theta) dz.$$

Factor Analysis

Factor Analysis




Factor Analysis: Generative Model




$$p(z) = \mathcal{N}(z|\mathbf{0}, \mathbf{I})$$

- Can be any Gaussian (see Murphy, book 1, section 20.2)



Factor Analysis: Generative Model




$$p(\mathbf{z}) = \mathcal{N}(\mathbf{z}|\mathbf{0}, \mathbf{I})$$



$$p(\mathbf{x}|\mathbf{z}) = \mathcal{N}(\mathbf{x}|\mathbf{W}\mathbf{z} + \boldsymbol{\mu}, \boldsymbol{\Psi})$$

- Can be any Gaussian (see Murphy, book 1, section 20.2)
- Linear Gaussian model
- $\mathbf{z} : D \times T$ latent. dim x samples (timepoints)
- $\mathbf{x} : N \times T$ obs. dim (neurons) x samples (timepoints)
- $\mathbf{W} : N \times D$ obs. dim. (neurons) x latent dim.
- $\boldsymbol{\Psi} : N \times N$ diagonal covariance matrix

Factor Analysis: Generative Model



$$p(z) = \mathcal{N}(z|\mathbf{0}, \mathbf{I})$$




$$p(x|z) = \mathcal{N}(x|\mathbf{W}z + \mu, \Psi)$$

$$p(x) = \int p(x|z)p(z)dz$$



$$p(x) = \mathcal{N}(x|\mu, \mathbf{W}\mathbf{W}^T + \Psi)$$

- Can be any Gaussian (see Murphy, book 1, section 20.2)
- Linear Gaussian model
- $z : D \times T$ latent. dim x samples (timepoints)
- $x : N \times T$ obs. dim (neurons) x samples (timepoints)
- $\mathbf{W} : N \times D$ obs. dim. (neurons) x latent dim.
- $\Psi : N \times N$ diagonal covariance matrix

Factor Analysis: Generative Model



$$p(z) = \mathcal{N}(z|\mathbf{0}, \mathbf{I})$$



$$p(x|z) = \mathcal{N}(x|\mathbf{W}z + \mu, \Psi)$$

$$p(x) = \int p(x|z)p(z)dz$$

$$p(x) = \mathcal{N}(x|\mu, \mathbf{W}\mathbf{W}^T + \Psi)$$




Low Rank + Noise

Shared and Unique

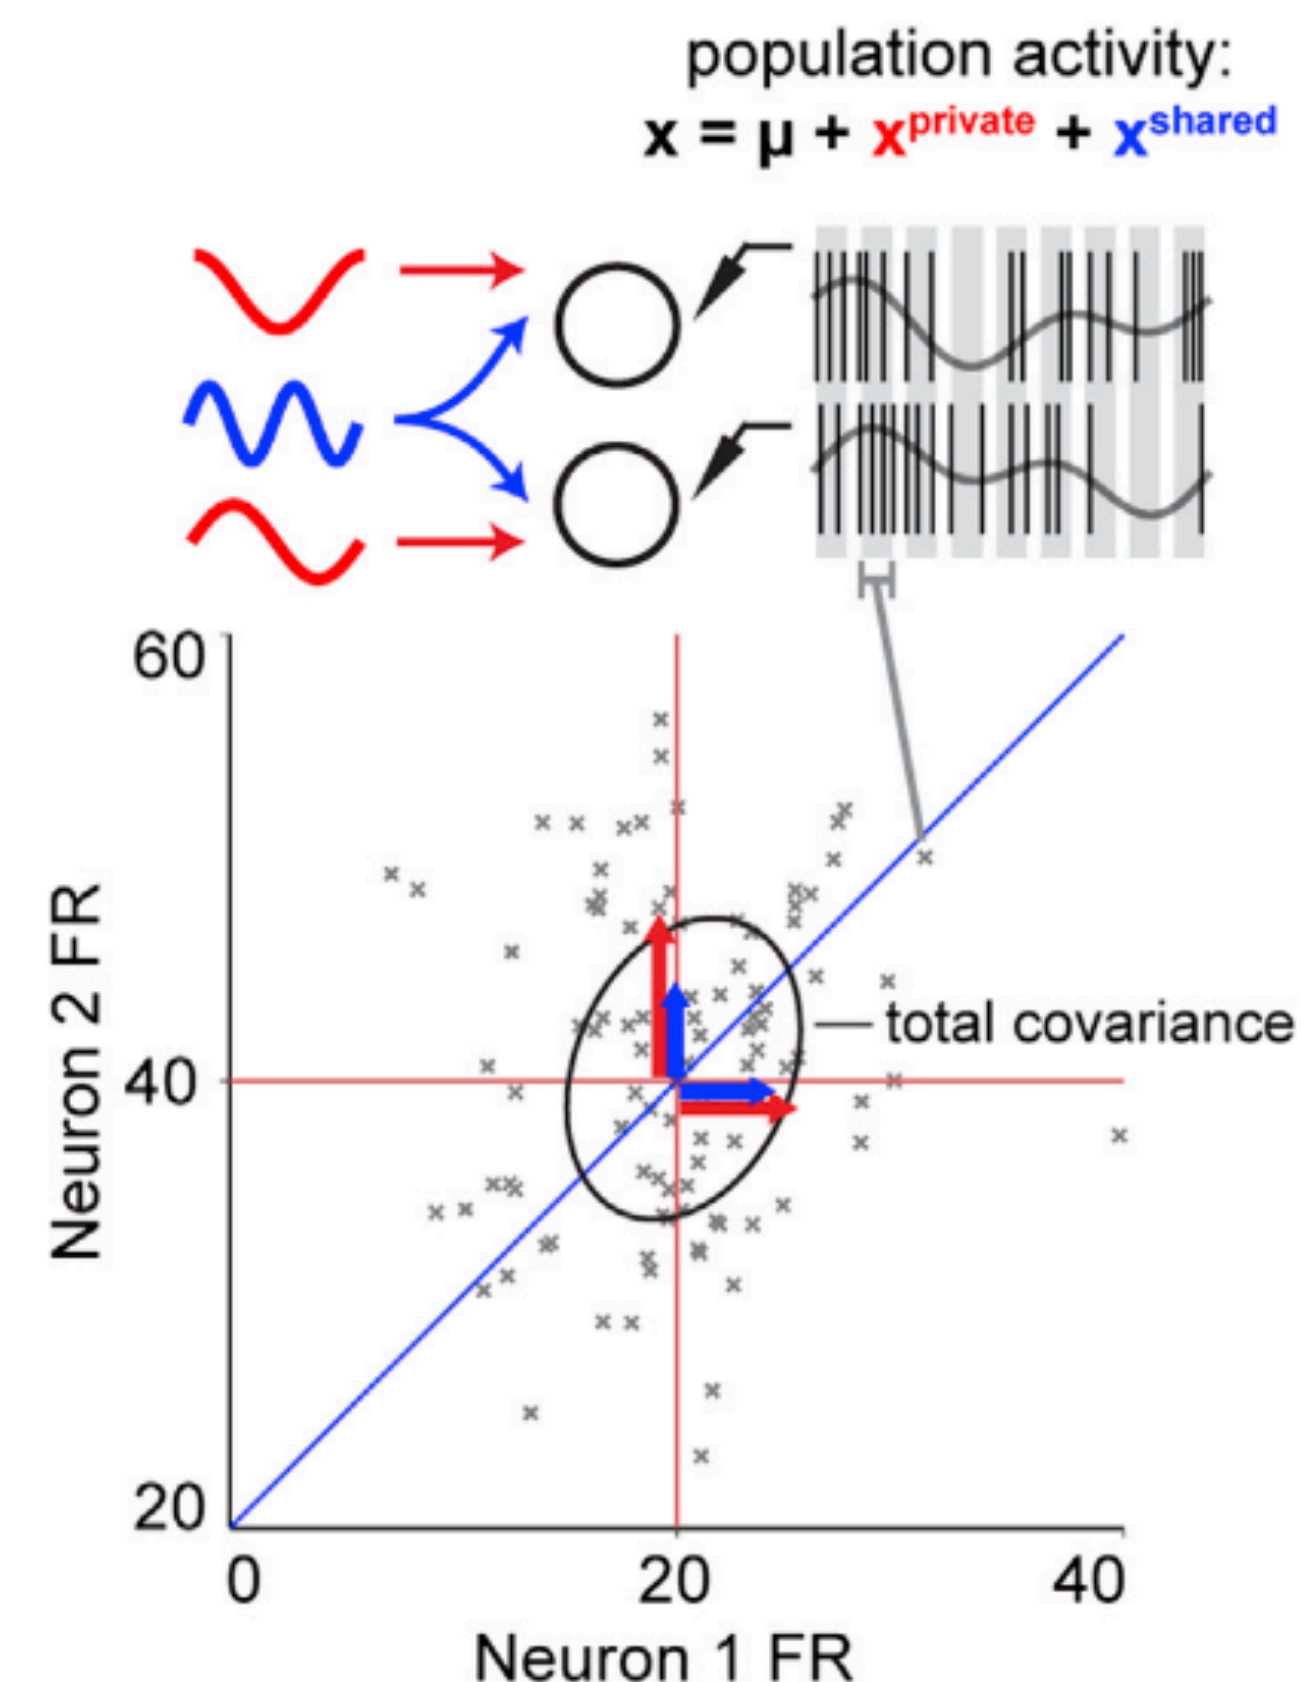
- Can be any Gaussian (see Murphy, book 1, section 20.2)
- Linear Gaussian model
- $z : D \times T$ latent. dim x samples (timepoints)
- $x : N \times T$ obs. dim (neurons) x samples (timepoints)
- $\mathbf{W} : N \times D$ obs. dim. (neurons) x latent dim.
- $\Psi : D \times D$ diagonal covariance matrix

Factor Analysis: Shared and Unique Example

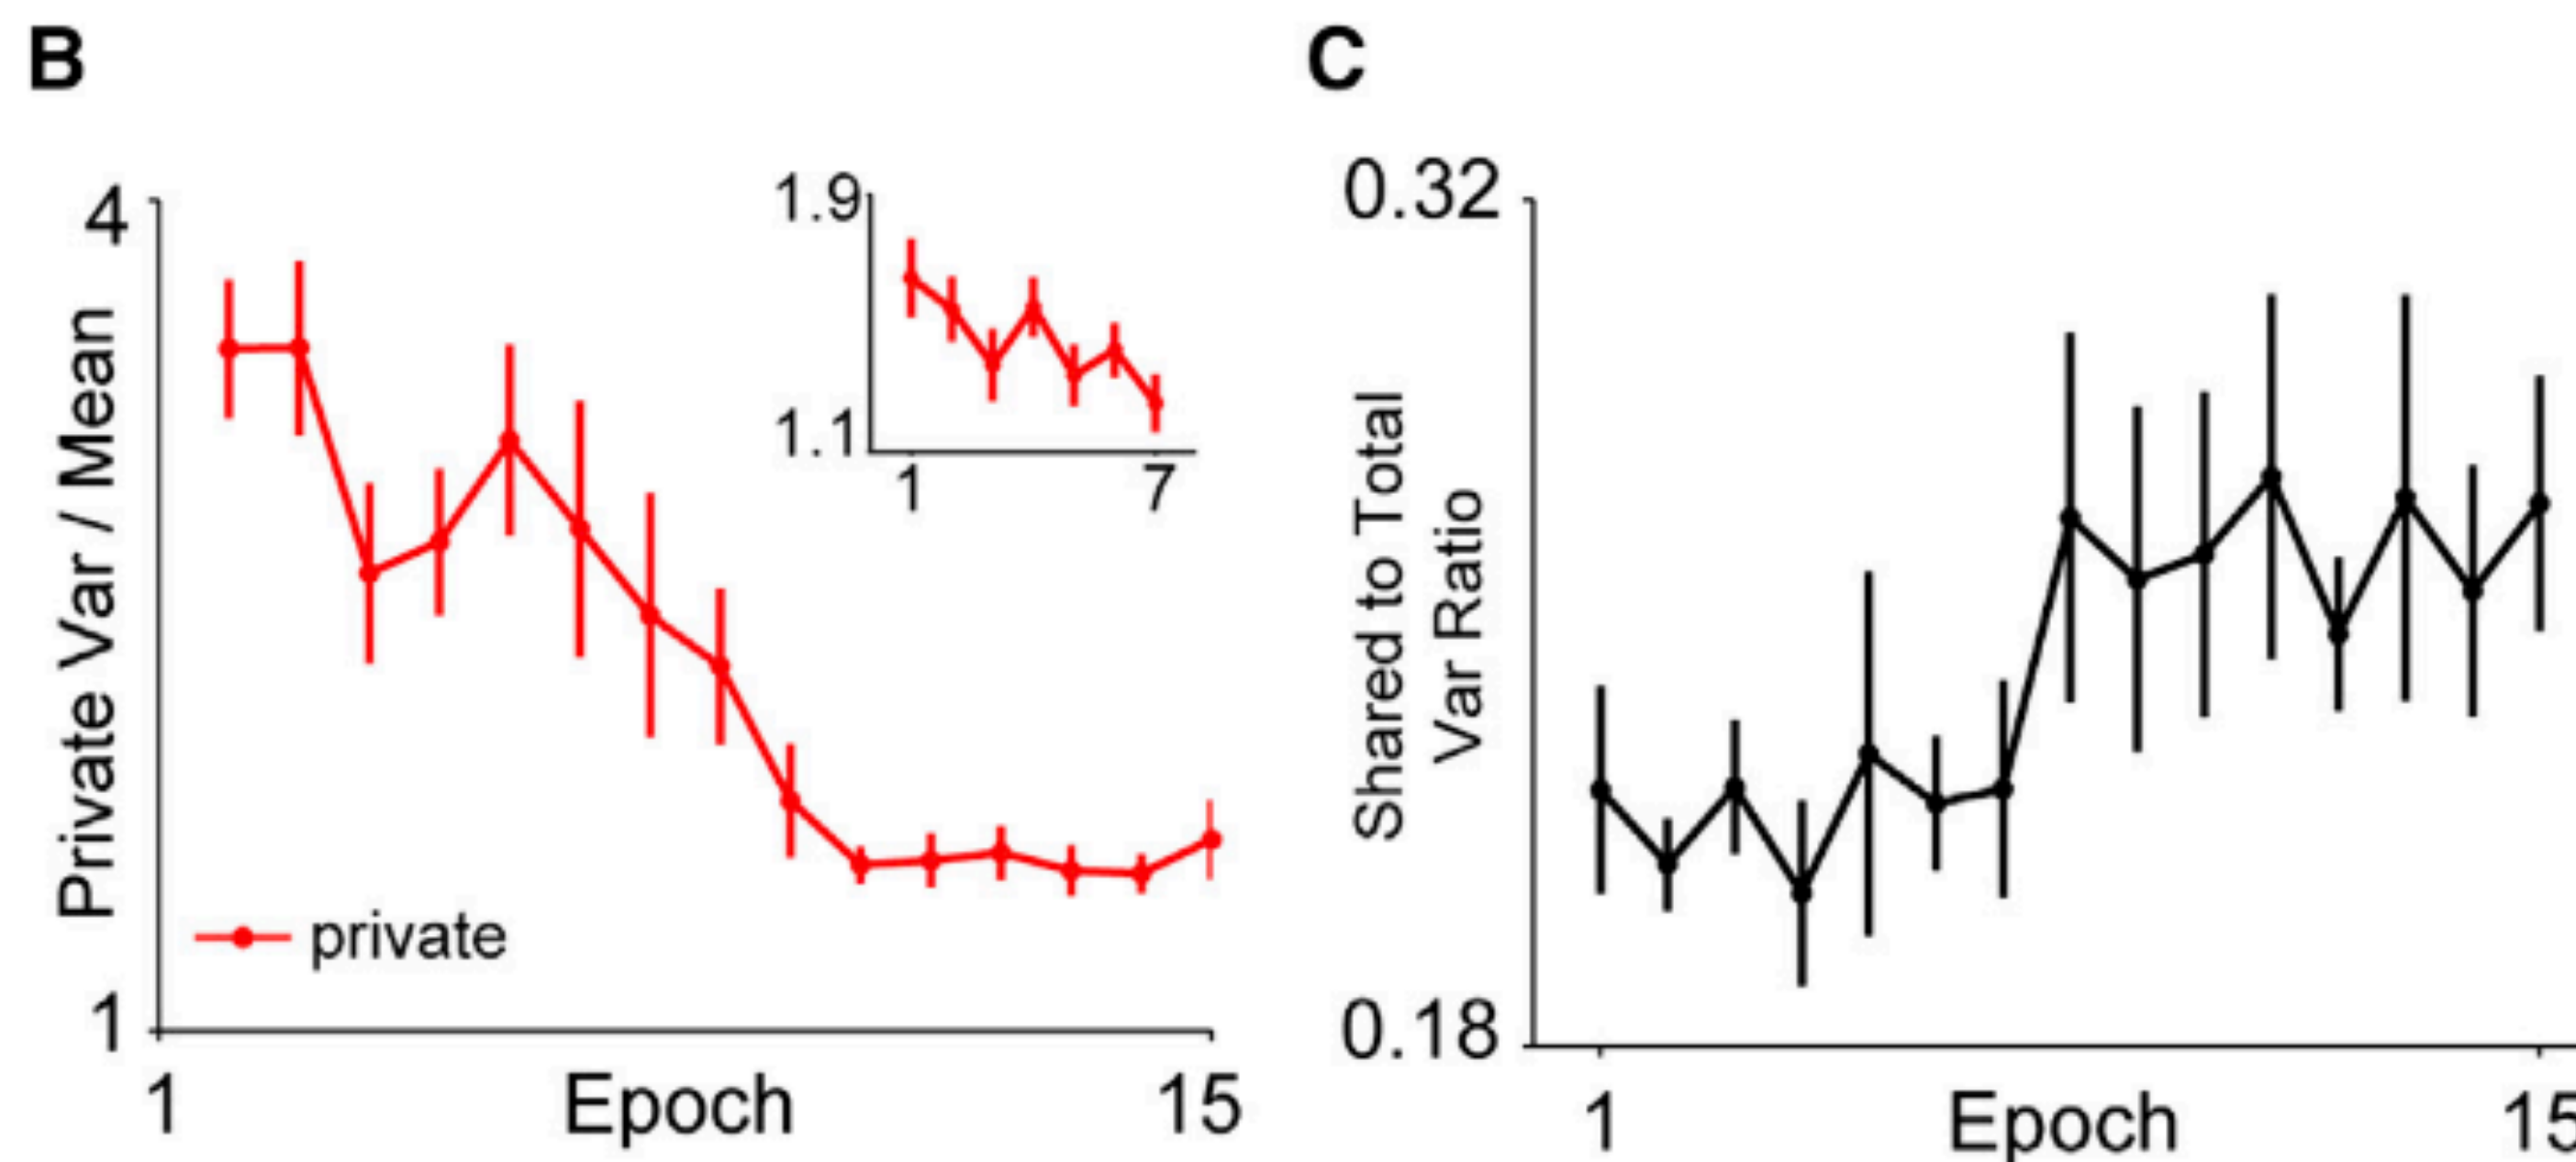
D



G




Factor Analysis: Shared and Unique Example



FA vs. Probabilistic PCA vs. PCA




$$p(z) = \mathcal{N}(z|\mathbf{0}, \mathbf{I})$$




$$p(x|z) = \mathcal{N}(x|\mathbf{W}z + \mu, \Psi)$$

FA vs. Probabilistic PCA vs. PCA



- Probabilistic PCA is Factor Analysis where Ψ is the identity matrix (all observations have the same independent noise)

FA vs. Probabilistic PCA vs. PCA




- Probabilistic PCA is Factor Analysis where Ψ is the identity matrix (all observations have the same independent noise)
- PPCA when $\Psi \rightarrow 0$ becomes PCA

FA vs. PCA: An Example



$$p(z) = \mathcal{N}(z|\mathbf{0}, \mathbf{I})$$




$$p(x|z) = \mathcal{N}(x|\mathbf{W}z + \mu, \Psi)$$

$$\mathbf{W} = \begin{bmatrix} 1 \\ 1 \end{bmatrix}$$

$$\Psi = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$


FA vs. PCA: An Example



$$W = \begin{bmatrix} 1 \\ 1 \end{bmatrix} \quad \Psi = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

$$\text{cov}(X) = WW^T + \Psi = \begin{bmatrix} 1 & 1 \\ 1 & 2 \end{bmatrix}$$

FA vs. PCA: An Example



$$W = \begin{bmatrix} 1 \\ 1 \end{bmatrix} \quad \Psi = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

$$\text{cov}(X) = WW^T + \Psi = \begin{bmatrix} 1 & 1 \\ 1 & 2 \end{bmatrix}$$

- PCA would give a top eigenvector primarily lying along the first dimension

Factor Analysis: Inferring the latents

$$\begin{aligned} p(z|x) &\propto p(x|z)p(z) \\ &= \mathcal{N}(x|Wz, \Psi) \cdot \mathcal{N}(z|0, I) \end{aligned}$$

Factor Analysis: Inferring the latents

$$p(z|x) \propto p(x|z)p(z)$$

$$= \mathcal{N}(x|Wz, \Psi) \cdot \mathcal{N}(z|0, I)$$

$$\vdots$$

$$= \mathcal{N}(\Lambda W^T \Psi^{-1} x, \Lambda) \quad \text{where} \quad \Lambda = (W^T \Psi^{-1} W + I)^{-1}$$

Factor Analysis: Inferring the latents

$$p(z | x) \propto p(x | z)p(z)$$

$$= \mathcal{N}(x | Wz, \Psi) \cdot \mathcal{N}(z | 0, I)$$

$$\vdots$$

$$= \mathcal{N}(\Lambda W^T \Psi^{-1} x, \Lambda) \quad \text{where} \quad \Lambda = (W^T \Psi^{-1} W + I)^{-1}$$

- When inferring the latent, the components of x are downweighted in proportion to their amount of independent noise (value in Ψ).

EM for Factor Analysis

- E step: Estimate the posterior, $p(z|x)$, given set parameters
- M step: Estimate the parameters, $[W, \Psi]$, given the expectations of the latents

Why probabilistic models, versus PCA?

- Allows having more sophisticated, and more accurate models
 - Different noise models (FA vs PPCA), mixture of factor analyzers, etc...
- Principled
- Better for missing data, or streaming data
- FA won't be as dependent on scaling

- However, PCA has simpler understanding in terms of variance and orthogonality, and is faster to run!
- Important to check that scientific results are robust across methods

Gaussian Processes

Gaussian Processes

Now consider a function $f : \mathcal{X} \rightarrow \mathbb{R}$ evaluated at a set of inputs, $\mathbf{X} = \{x_n \in \mathcal{X}\}_{n=1}^N$. Let $\mathbf{f}_X = [f(x_1), \dots, f(x_N)]$ be the set of unknown function values at these points.

Gaussian Processes

Now consider a function $f : \mathcal{X} \rightarrow \mathbb{R}$

evaluated at a set of inputs, $\mathbf{X} = \{x_n \in \mathcal{X}\}_{n=1}^N$. Let $\mathbf{f}_X = [f(x_1), \dots, f(x_N)]$ be the set of unknown function values at these points. If \mathbf{f}_X is jointly Gaussian for any set of $N \geq 1$ points, then we say that $f : \mathcal{X} \rightarrow \mathbb{R}$ is a **Gaussian process**.

Gaussian Processes

Now consider a function $f : \mathcal{X} \rightarrow \mathbb{R}$ evaluated at a set of inputs, $\mathbf{X} = \{x_n \in \mathcal{X}\}_{n=1}^N$. Let $\mathbf{f}_X = [f(x_1), \dots, f(x_N)]$ be the set of unknown function values at these points. If \mathbf{f}_X is jointly Gaussian for any set of $N \geq 1$ points, then we say that $f : \mathcal{X} \rightarrow \mathbb{R}$ is a **Gaussian process**. Such a process is defined by its **mean function** $m(x) \in \mathbb{R}$ and a **covariance function**, $\mathcal{K}(x, x') \geq 0$, which is any positive definite **Mercer kernel**

Gaussian Processes


Now consider a function $f : \mathcal{X} \rightarrow \mathbb{R}$ evaluated at a set of inputs, $\mathbf{X} = \{\mathbf{x}_n \in \mathcal{X}\}_{n=1}^N$. Let $\mathbf{f}_X = [f(\mathbf{x}_1), \dots, f(\mathbf{x}_N)]$ be the set of unknown function values at these points. If \mathbf{f}_X is jointly Gaussian for any set of $N \geq 1$ points, then we say that $f : \mathcal{X} \rightarrow \mathbb{R}$ is a **Gaussian process**. Such a process is defined by its **mean function** $m(\mathbf{x}) \in \mathbb{R}$ and a **covariance function**, $\mathcal{K}(\mathbf{x}, \mathbf{x}') \geq 0$, which is any positive definite **Mercer kernel**

- Example Kernel (“Radial Basis Function”): $\mathcal{K}(\mathbf{x}, \mathbf{x}'; \ell) = \exp\left(-\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{2\ell^2}\right)$

Gaussian Processes

Now consider a function $f : \mathcal{X} \rightarrow \mathbb{R}$ evaluated at a set of inputs, $\mathbf{X} = \{x_n \in \mathcal{X}\}_{n=1}^N$. Let $\mathbf{f}_X = [f(x_1), \dots, f(x_N)]$ be the set of unknown function values at these points. If \mathbf{f}_X is jointly Gaussian for any set of $N \geq 1$ points, then we say that $f : \mathcal{X} \rightarrow \mathbb{R}$ is a **Gaussian process**. Such a process is defined by its **mean function** $m(x) \in \mathbb{R}$ and a **covariance function**, $\mathcal{K}(x, x') \geq 0$, which is any positive definite **Mercer kernel**



- Example Kernel (“Radial Basis Function”): $\mathcal{K}(x, x'; \ell) = \exp\left(-\frac{\|x - x'\|^2}{2\ell^2}\right)$




Gaussian Processes

Now consider a function $f : \mathcal{X} \rightarrow \mathbb{R}$ evaluated at a set of inputs, $\mathbf{X} = \{x_n \in \mathcal{X}\}_{n=1}^N$. Let $\mathbf{f}_X = [f(x_1), \dots, f(x_N)]$ be the set of unknown function values at these points. If \mathbf{f}_X is jointly Gaussian for any set of $N \geq 1$ points, then we say that $f : \mathcal{X} \rightarrow \mathbb{R}$ is a **Gaussian process**. Such a process is defined by its **mean function** $m(x) \in \mathbb{R}$ and a **covariance function**, $\mathcal{K}(x, x') \geq 0$, which is any positive definite **Mercer kernel**

- Example Kernel (“Radial Basis Function”): $\mathcal{K}(x, x'; \ell) = \exp\left(-\frac{\|x - x'\|^2}{2\ell^2}\right)$



Gaussian Processes - sampling from the prior



(a)

Figure 18.7: Left: some functions sampled from a GP prior with RBF kernel. Middle: some samples from a GP posterior, after conditioning on 5 noise-free observations. Right: some samples from a GP posterior, after conditioning on 5 noisy observations. The shaded area represents $\mathbb{E}[f(\mathbf{x})] \pm 2\sqrt{\mathbb{V}[f(\mathbf{x})]}$. Adapted from Figure 2.2 of [RW06]. Generated by [gpr_demo_noise_free.ipynb](#).

Gaussian Processes - Example kernels

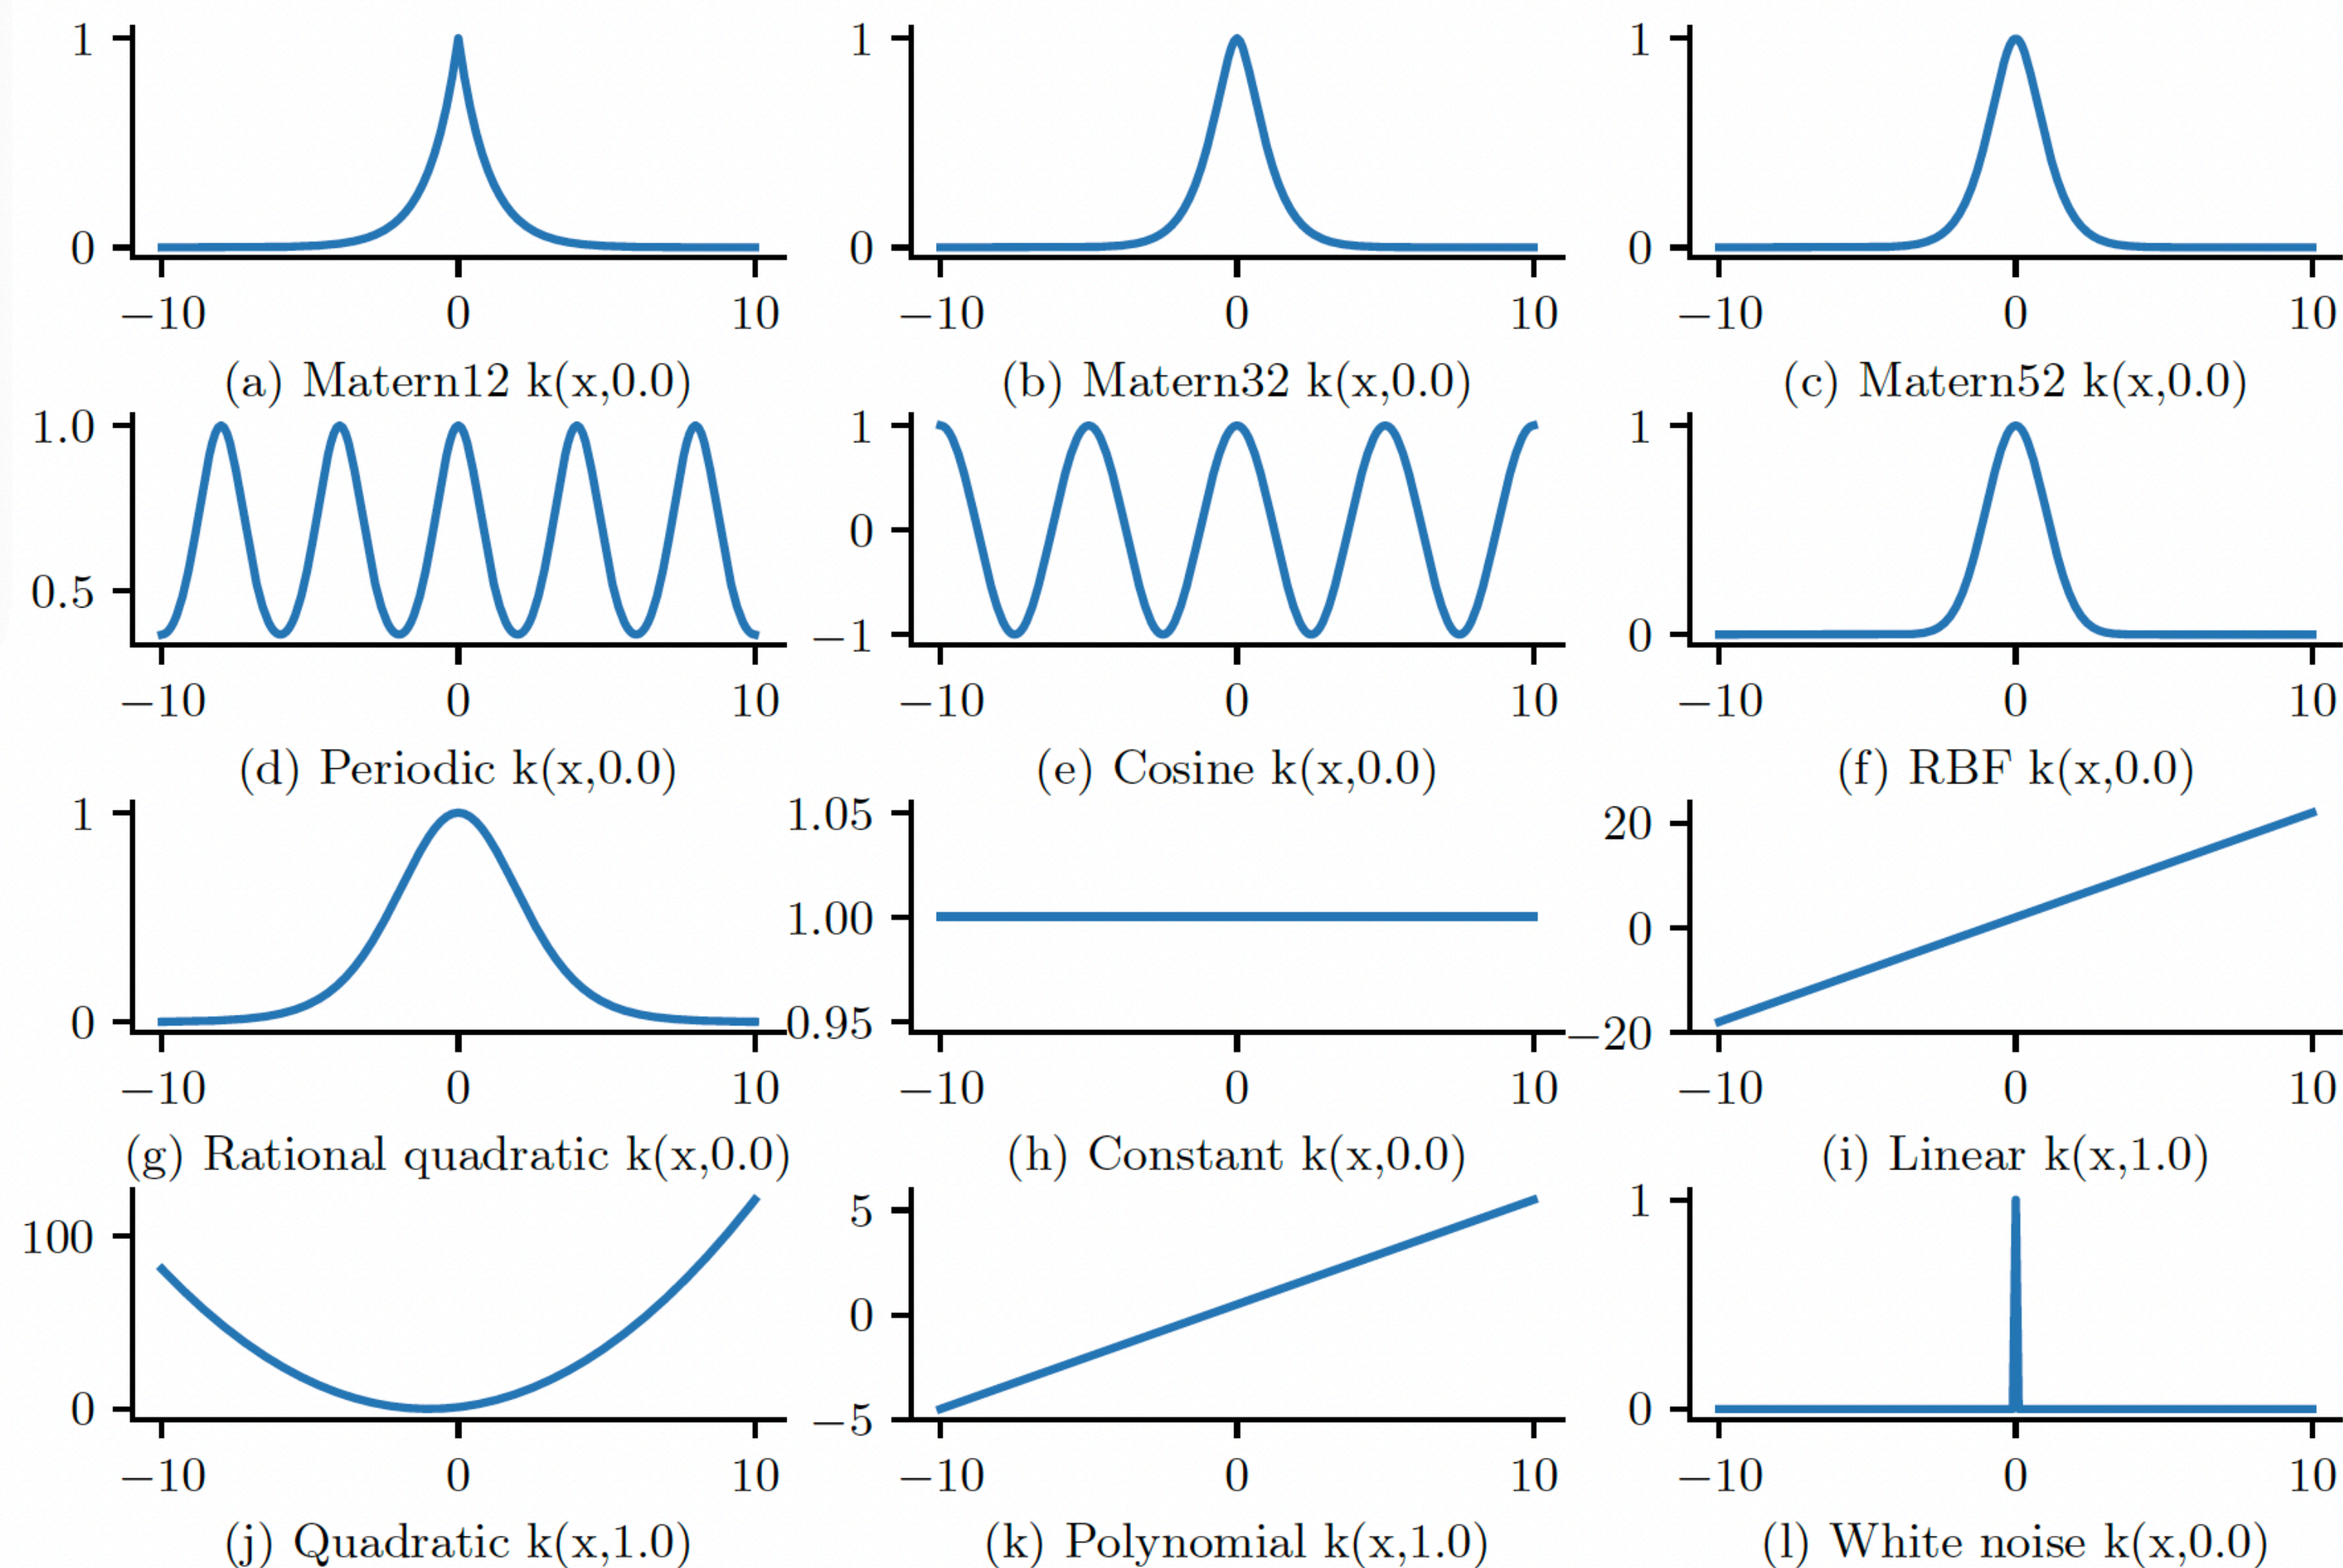



Figure 18.3: GP kernels evaluated at $k(x, 0)$ as a function of x . Generated by [gpKernelPlot.ipynb](#).

Gaussian Processes - estimating a posterior



(a)

Figure 18.7: Left: some functions sampled from a GP prior with RBF kernel. Middle: some samples from a GP posterior, after conditioning on 5 noise-free observations. Right: some samples from a GP posterior, after conditioning on 5 noisy observations. The shaded area represents $\mathbb{E}[f(\mathbf{x})] \pm 2\sqrt{\mathbb{V}[f(\mathbf{x})]}$. Adapted from Figure 2.2 of [RW06]. Generated by [gpr_demo_noise_free.ipynb](#).

Gaussian Processes - estimating a posterior

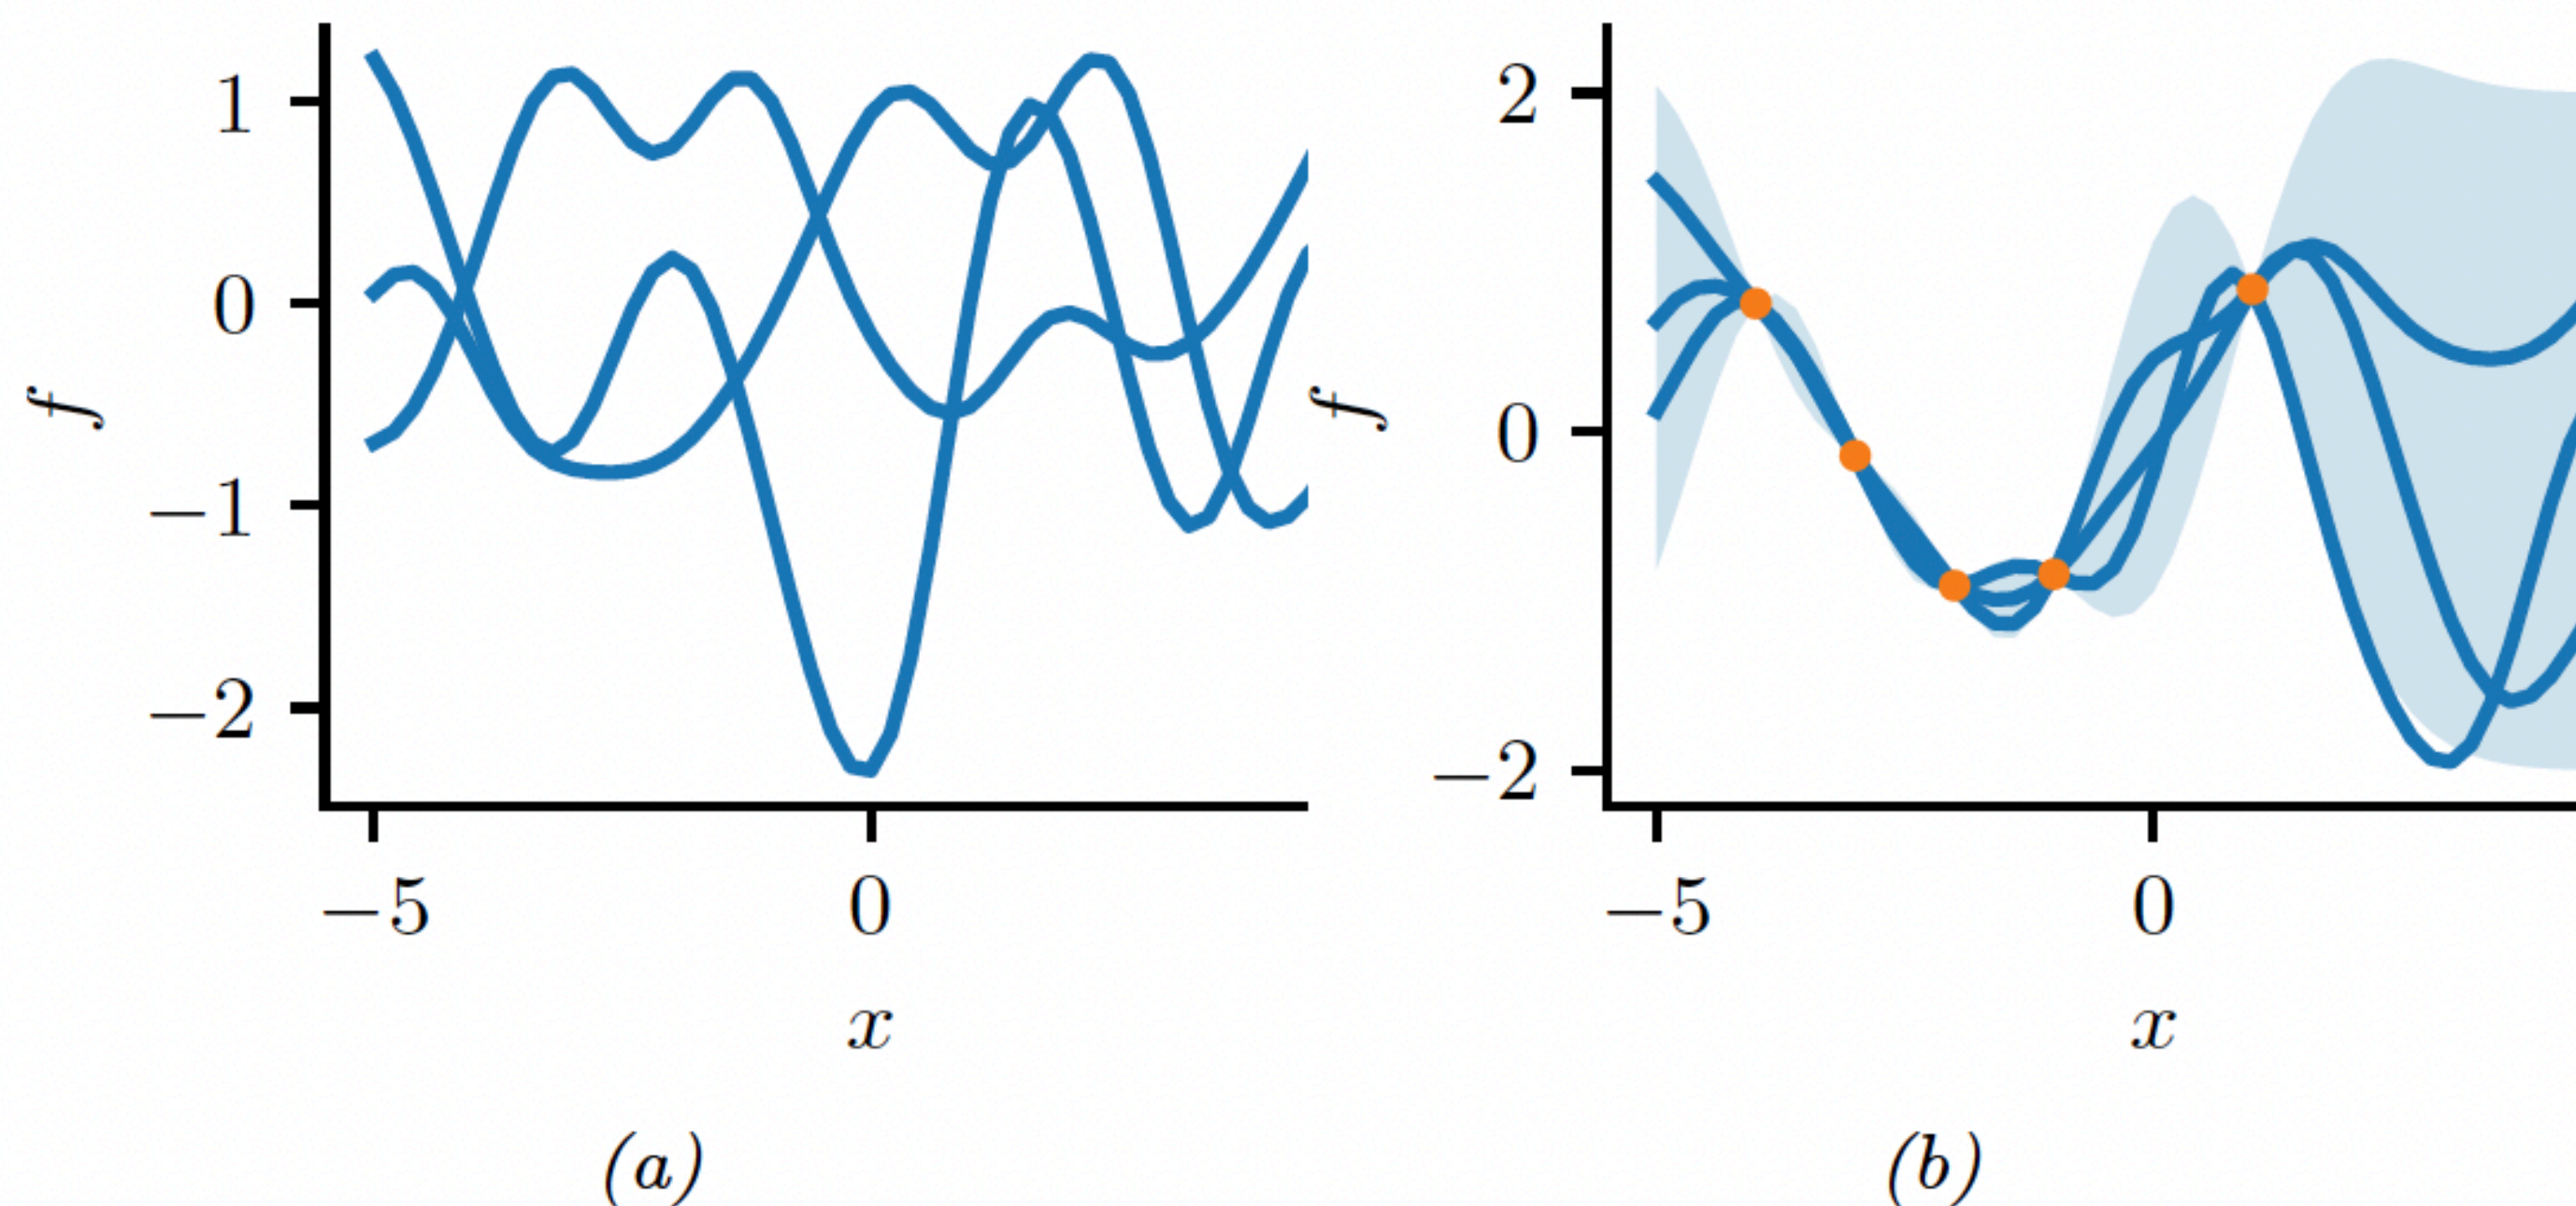


Figure 18.7: Left: some functions sampled from a GP prior with RBF kernel. Middle: some samples from a GP posterior, after conditioning on 5 noise-free observations. Right: some samples from a GP posterior, after conditioning on 5 noisy observations. The shaded area represents $\mathbb{E}[f(\mathbf{x})] \pm 2\sqrt{\mathbb{V}[f(\mathbf{x})]}$. Adapted from Figure 2.2 of [RW06]. Generated by [gpr_demo_noise_free.ipynb](#).

Gaussian Processes - estimating a posterior

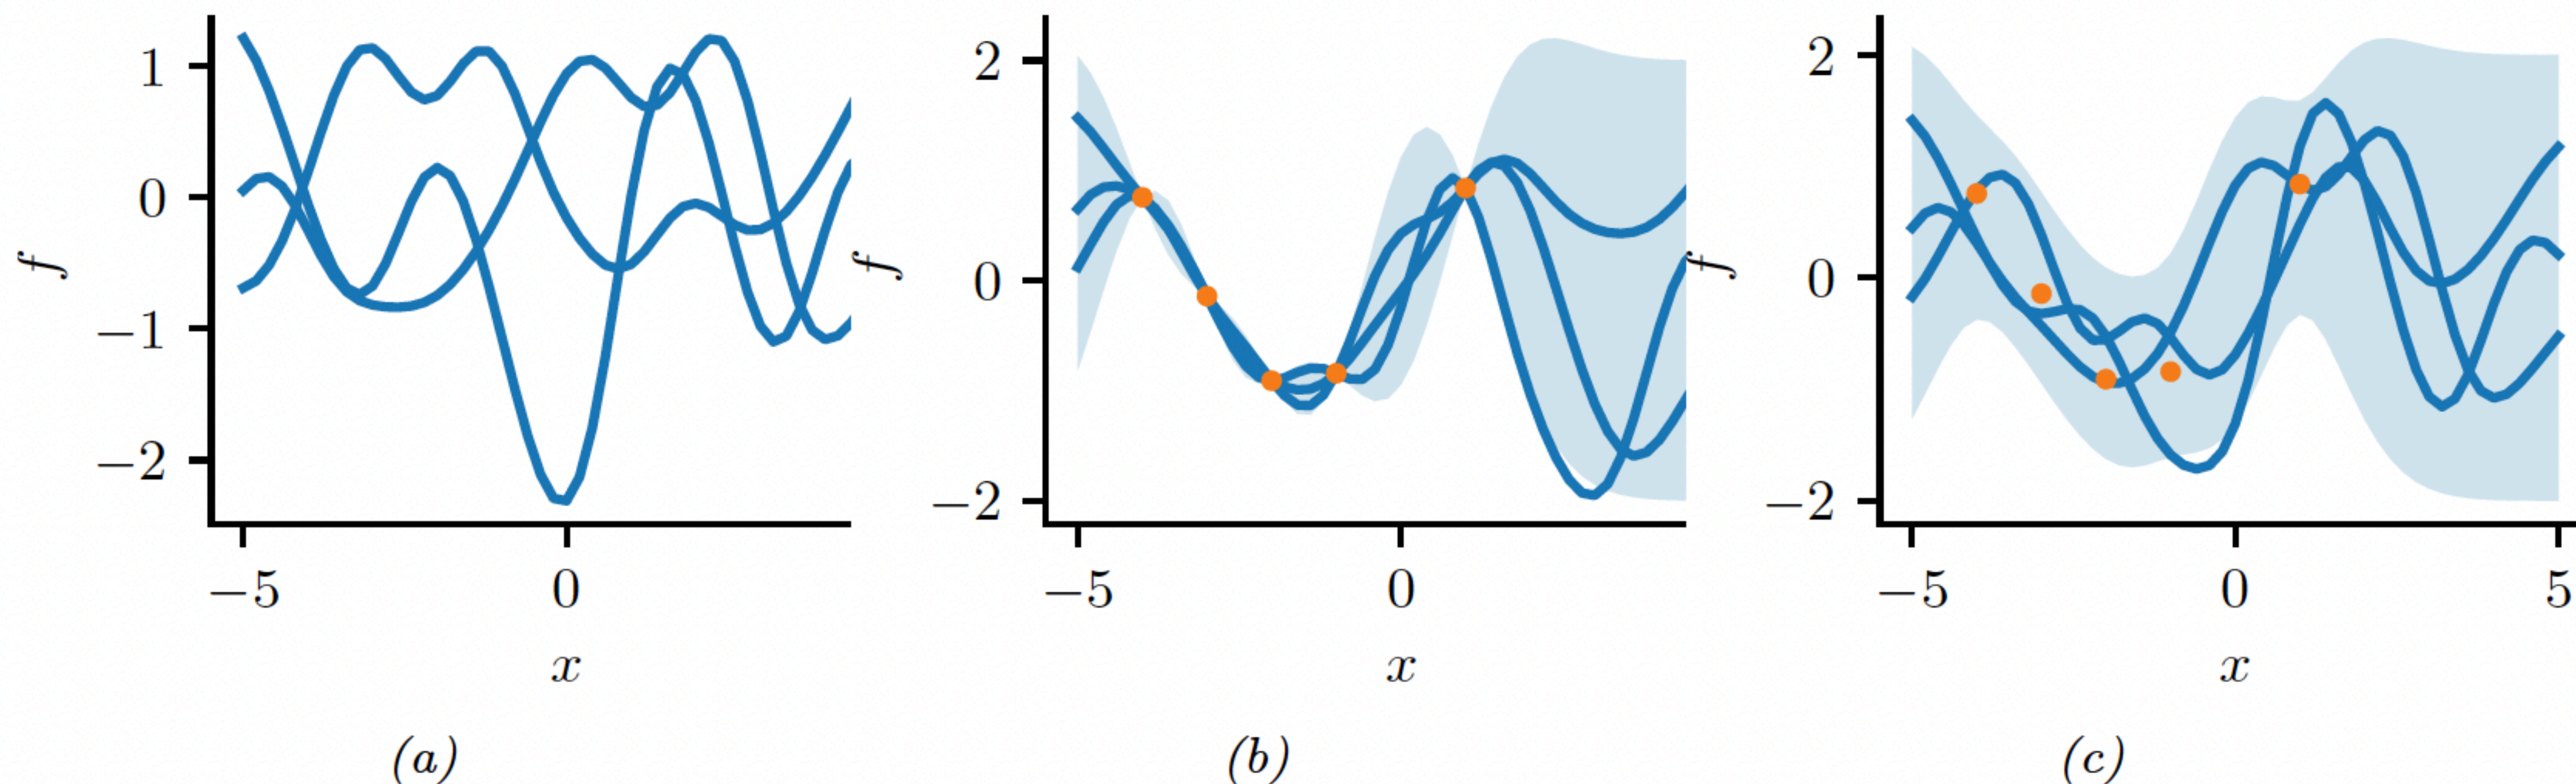


Figure 18.7: Left: some functions sampled from a GP prior with RBF kernel. Middle: some samples from a GP posterior, after conditioning on 5 noise-free observations. Right: some samples from a GP posterior, after conditioning on 5 noisy observations. The shaded area represents $\mathbb{E}[f(\mathbf{x})] \pm 2\sqrt{\mathbb{V}[f(\mathbf{x})]}$. Adapted from Figure 2.2 of [RW06]. Generated by [gpr_demo_noise_free.ipynb](#).

Gaussian Process Factor Analysis (GPFA)

Gaussian Process Factor Analysis (GPFA)

- Automatically find smooth latent trajectory when inputting noisy spiking data

Gaussian Process Factor Analysis (GPFA)

$$\mathbf{y}_{:,t} \mid \mathbf{x}_{:,t} \sim \mathcal{N}(C\mathbf{x}_{:,t} + \mathbf{d}, R),$$

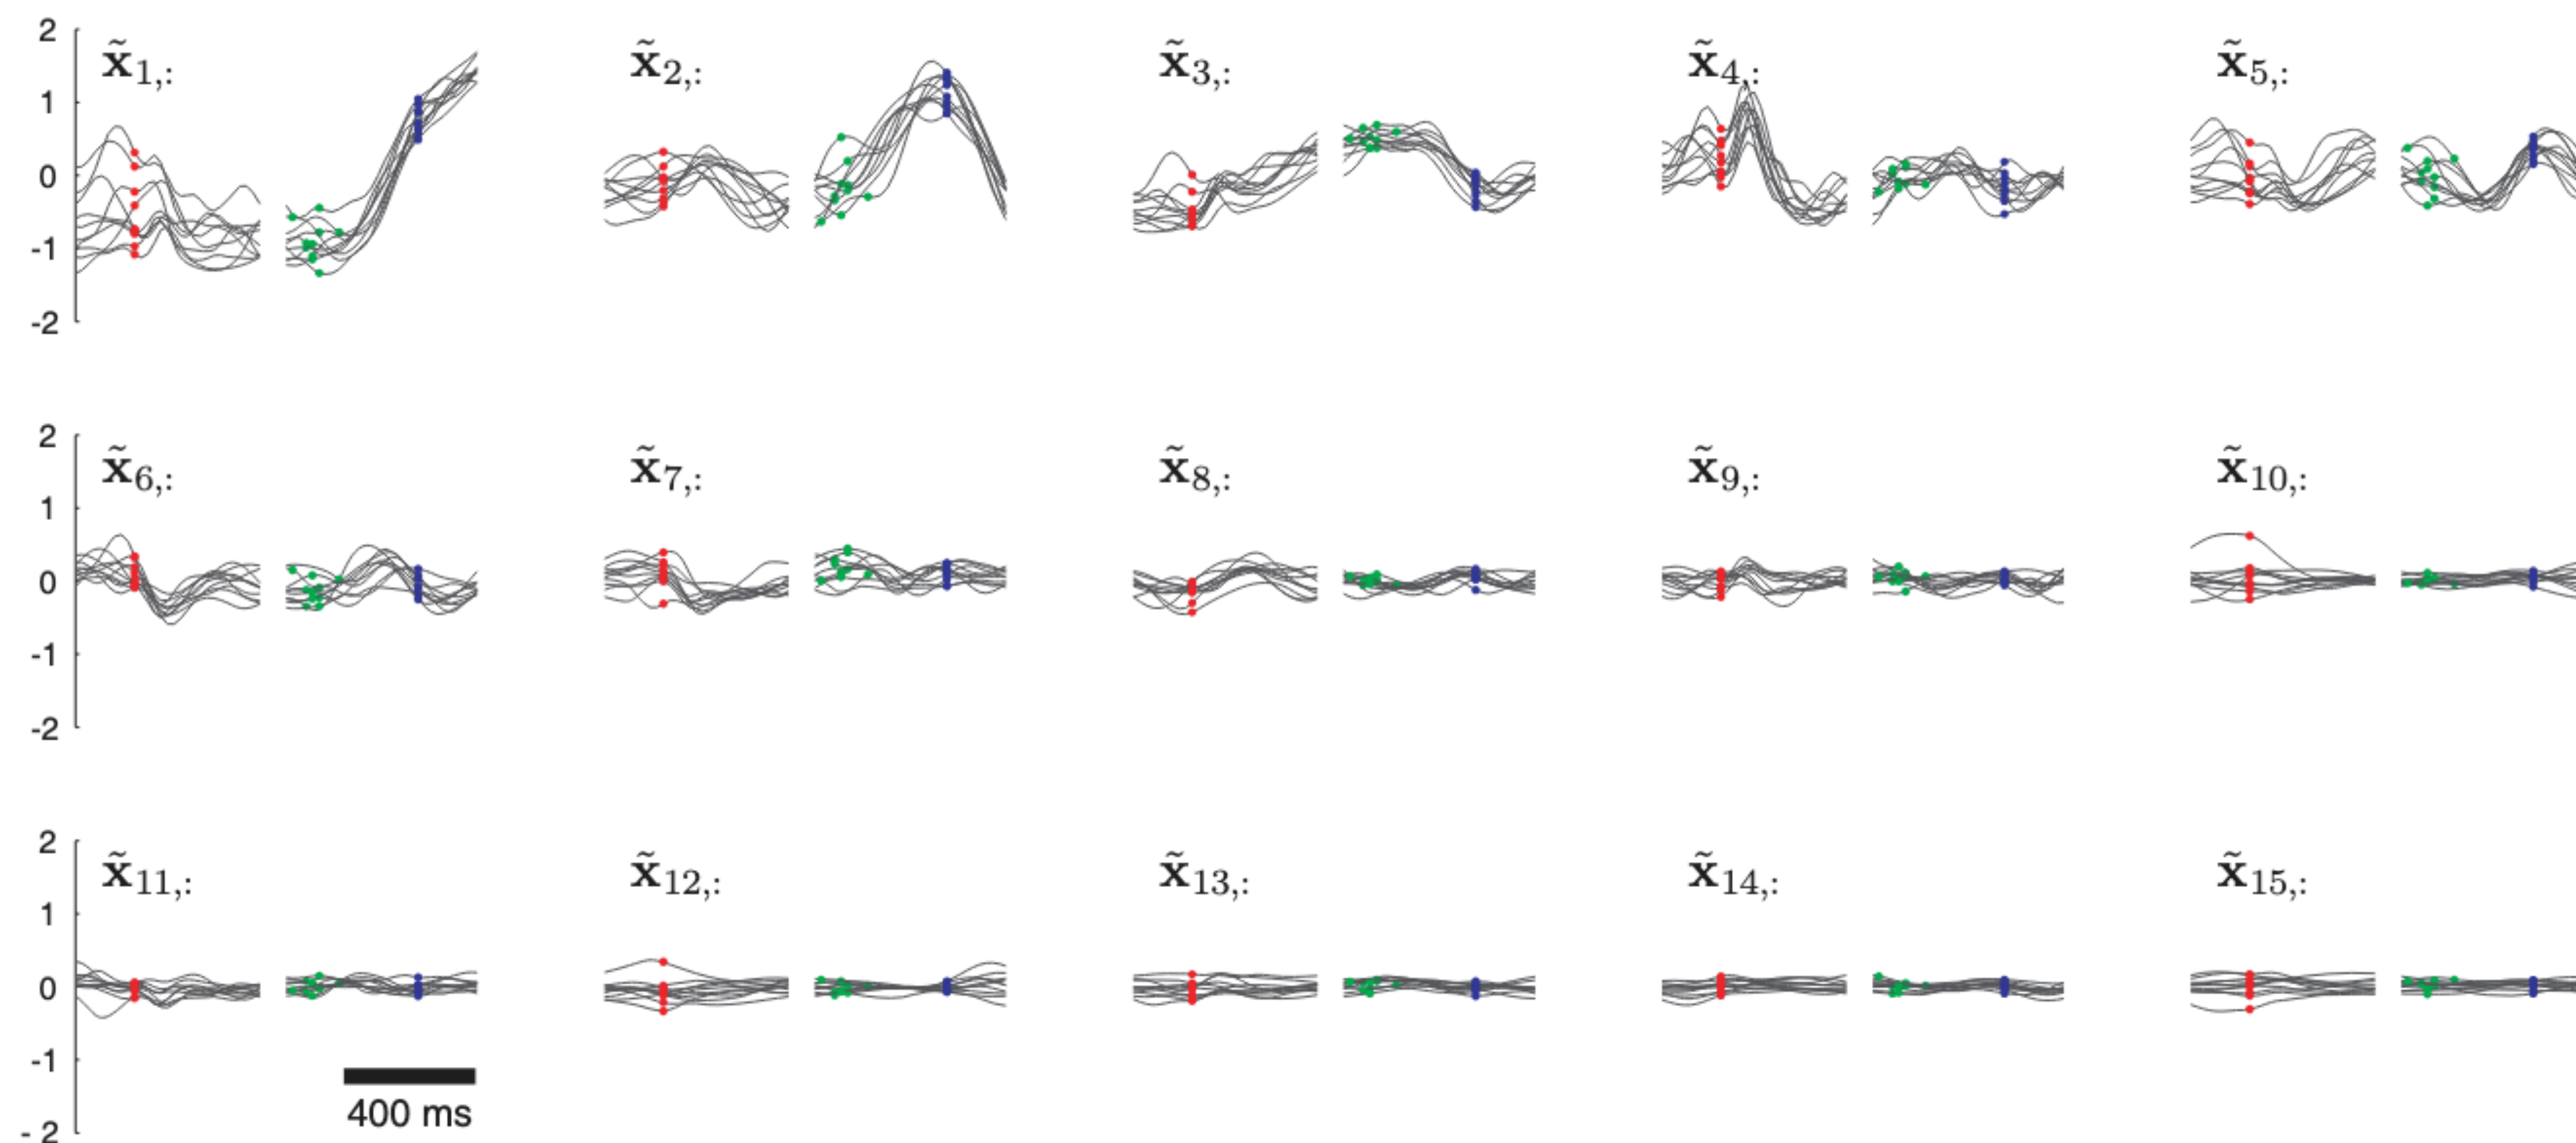
Gaussian Process Factor Analysis (GPFA)

$$\mathbf{y}_{:,t} \mid \mathbf{x}_{:,t} \sim \mathcal{N}(C\mathbf{x}_{:,t} + \mathbf{d}, R),$$

$$\mathbf{x}_{i,:} \sim \mathcal{N}(\mathbf{0}, K_i),$$

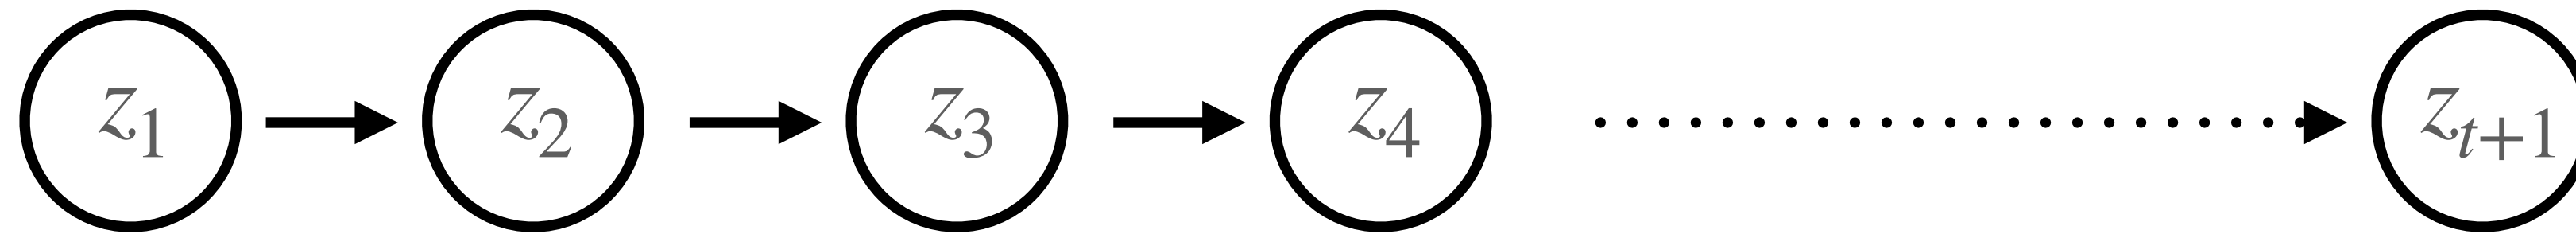
$$K_i(t_1, t_2) = \sigma_{f,i}^2 \cdot \exp\left(-\frac{(t_1 - t_2)^2}{2 \cdot \tau_i^2}\right) + \sigma_{n,i}^2 \cdot \delta_{t_1, t_2},$$

Gaussian Process Factor Analysis (GPFA)



HMMs

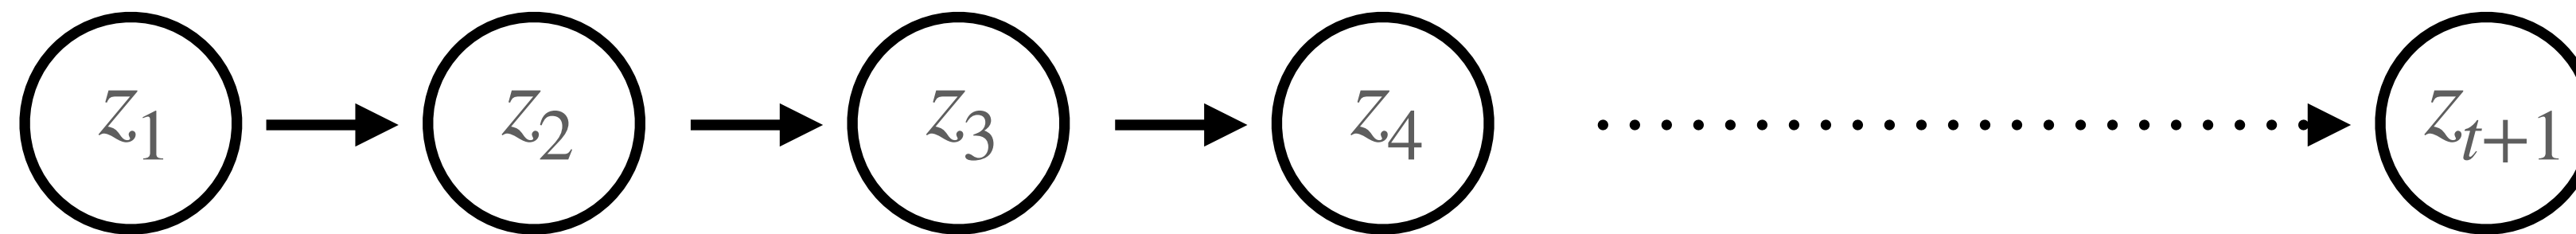
Markov Chains



- The current state only depends on the past state

$$P(z_{t+1} | z_1, z_2, \dots, z_t) = P(z_{t+1} | z_t)$$

Markov Chains



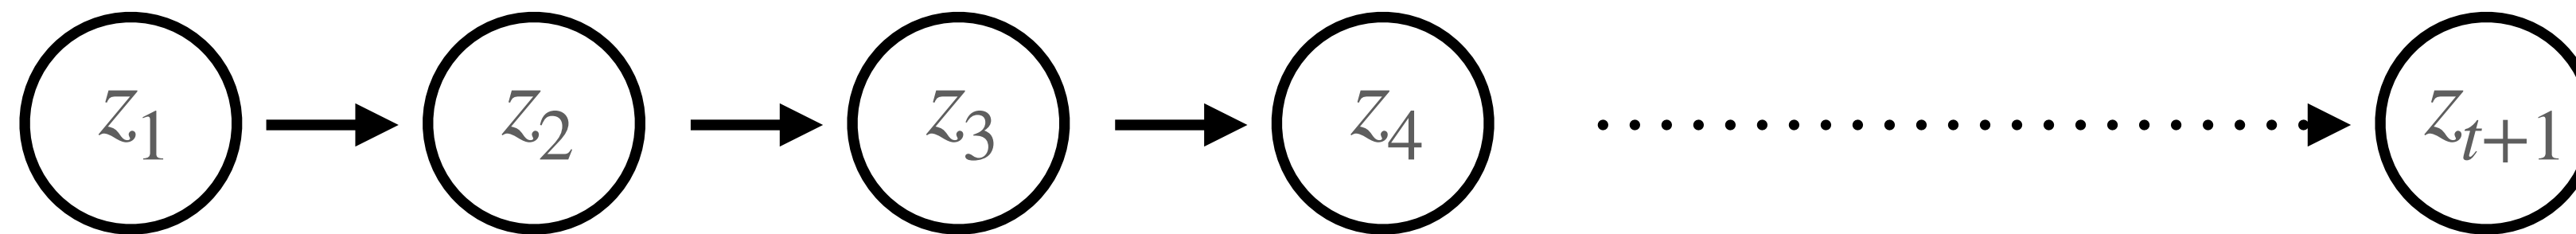
- The current state only depends on the past state

$$P(z_{t+1} | z_1, z_2, \dots, z_t) = P(z_{t+1} | z_t)$$

- We can use the rules of independence to calculate the total probability:

$$P(z_1, z_2, \dots, z_{t+1}) = P(z_1)P(z_2 | z_1) \dots P(z_t | z_{t-1})P(z_{t+1} | z_t)$$

Markov Chains



- The current state only depends on the past state

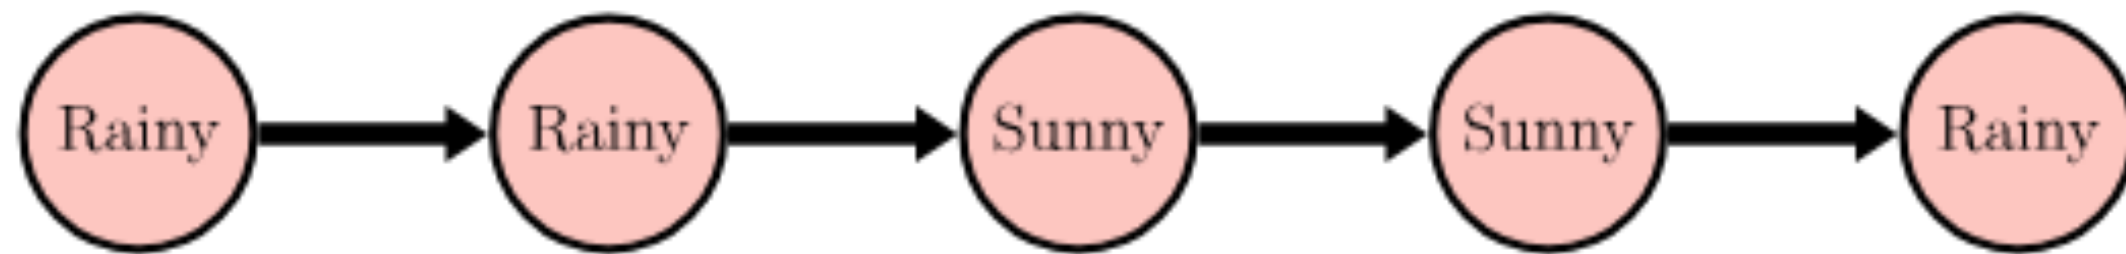
$$P(z_{t+1} | z_1, z_2, \dots, z_t) = P(z_{t+1} | z_t)$$

- We can use the rules of independence to calculate the total probability:


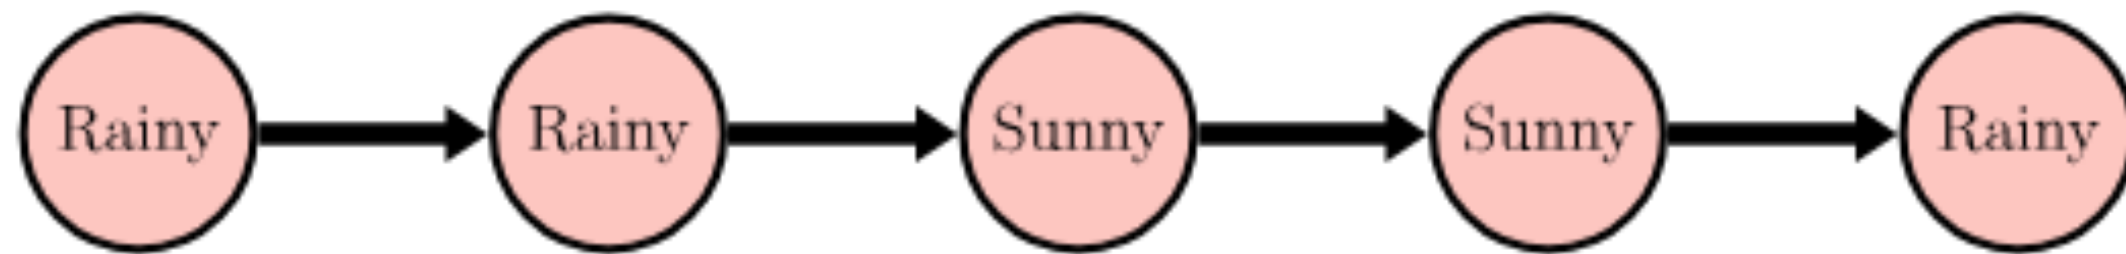
$$P(z_1, z_2, \dots, z_{t+1}) = P(z_1)P(z_2 | z_1) \dots P(z_t | z_{t-1})P(z_{t+1} | z_t)$$

$$P(z_{1:T}) = P(z_1) \prod_{t=2}^T p(z_t | z_{t-1})$$

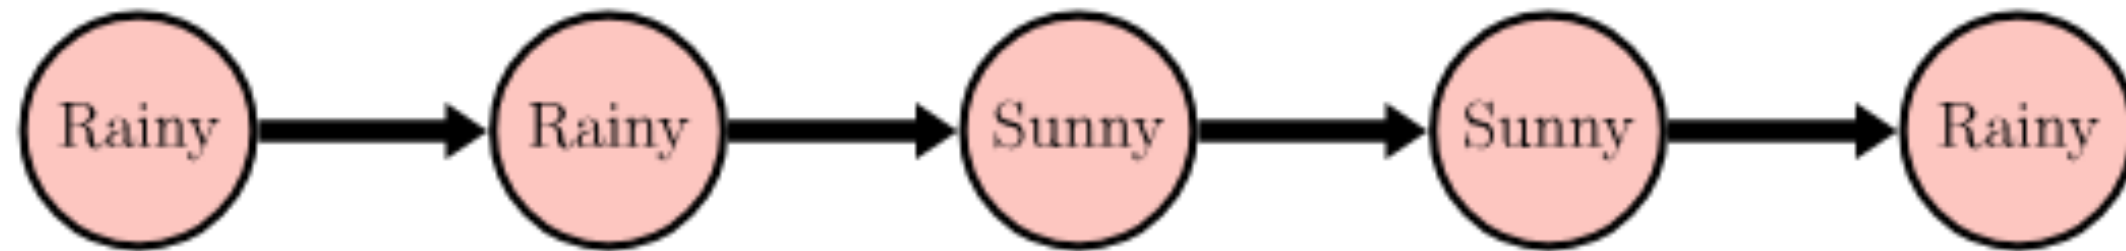
Markov Chains



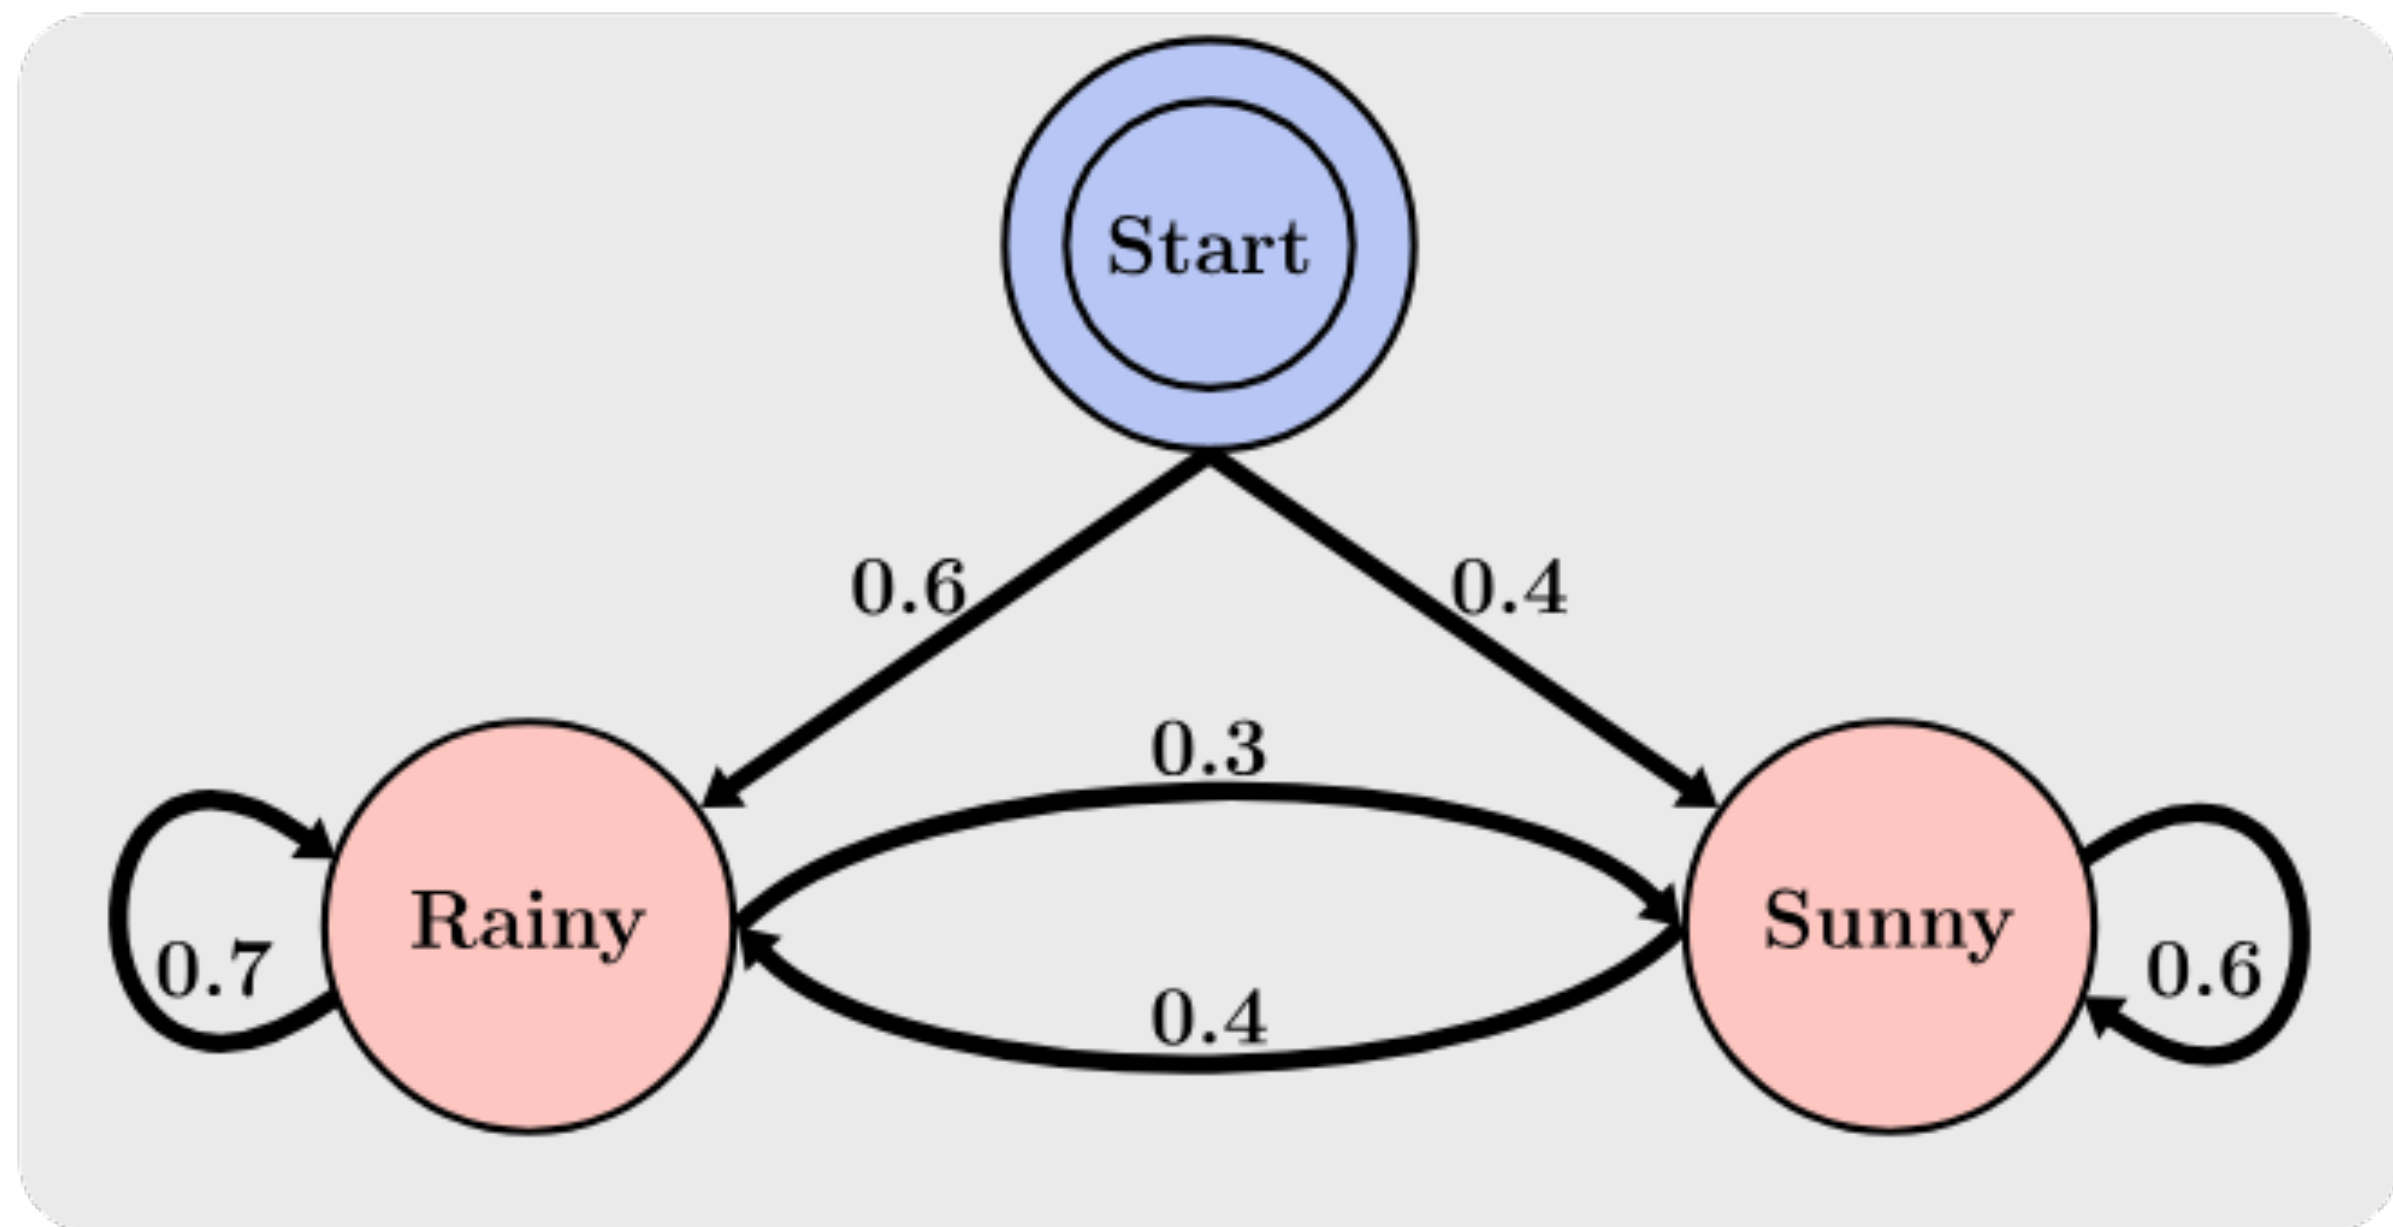
Markov Chains



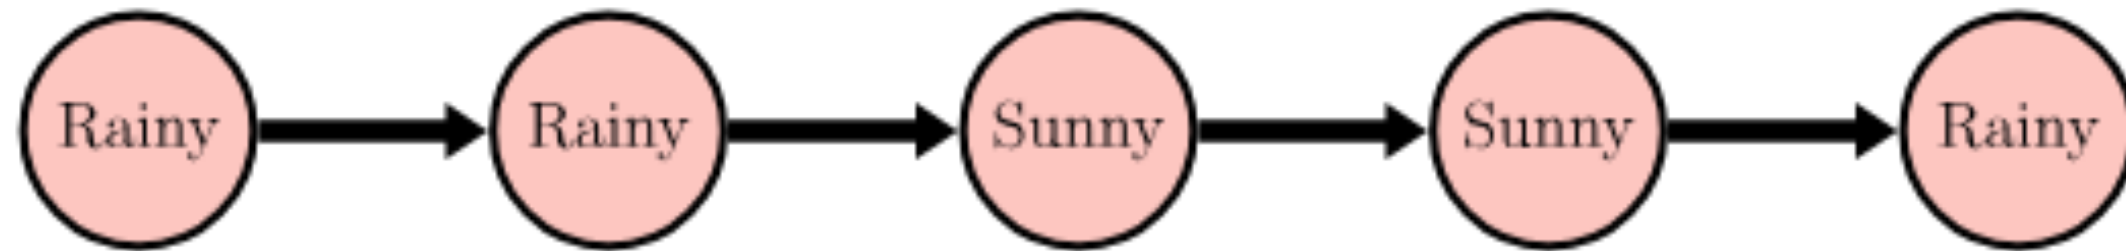
Markov Chains



$$P(z_1, z_2, \dots, z_{t+1}) = P(z_1)P(z_2 | z_1) \dots P(z_t | z_{t-1})P(z_{t+1} | z_t)$$



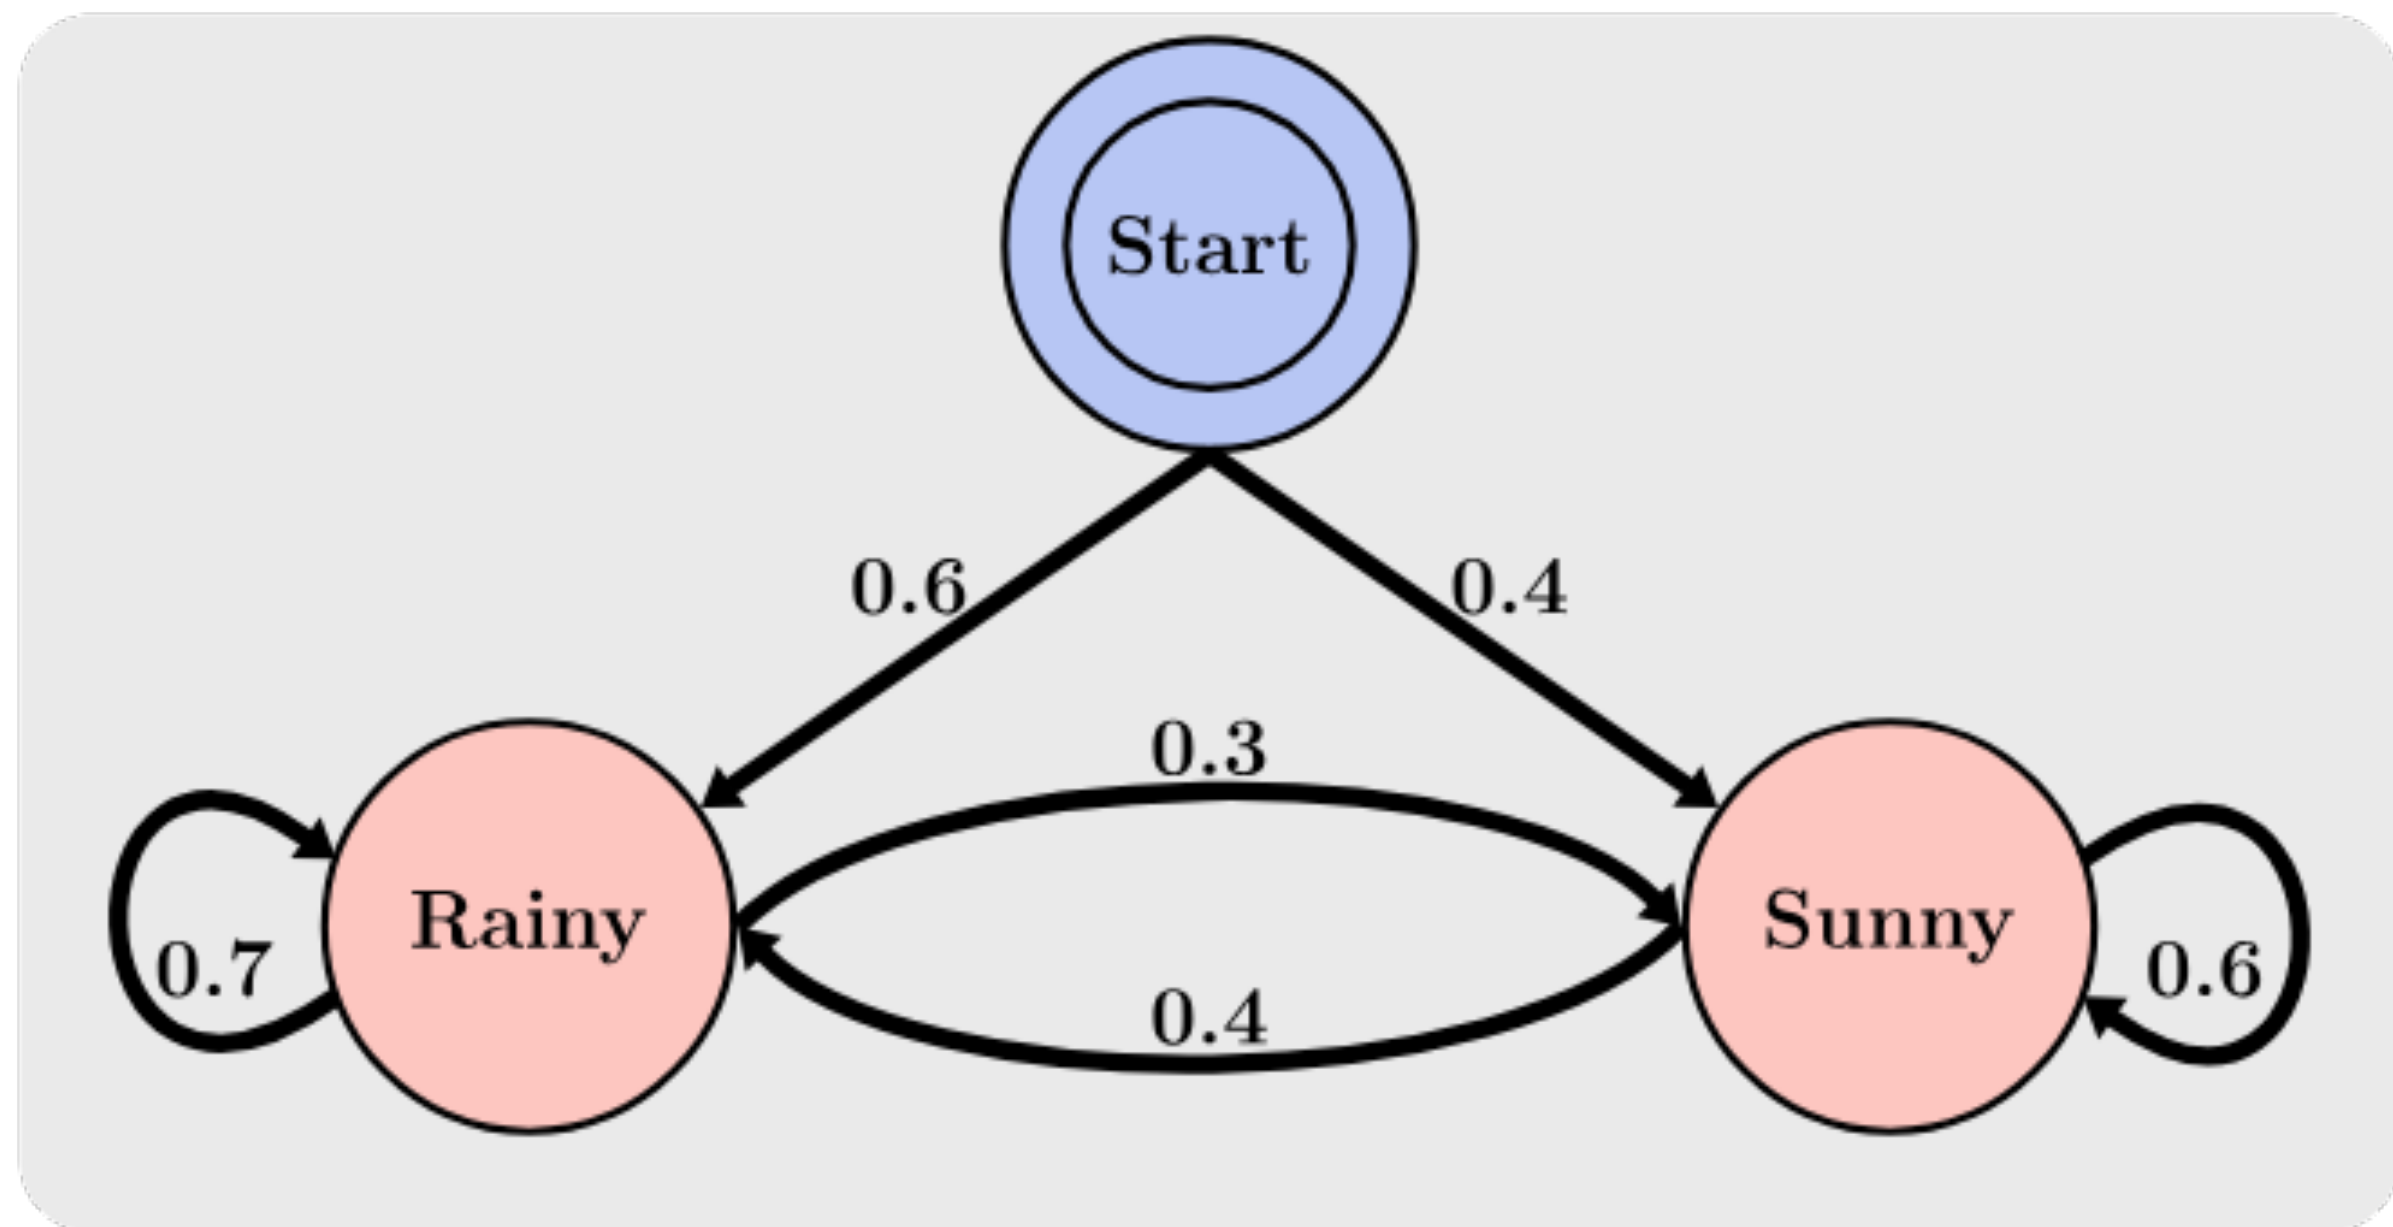
Markov Chains



$$P(z_1, z_2, \dots, z_{t+1}) = P(z_1)P(z_2 | z_1) \dots P(z_t | z_{t-1})P(z_{t+1} | z_t)$$

Initial Conditions

$$P(z_1 = R) = 0.6, \quad P(z_1 = S) = 0.4$$



Markov Chains



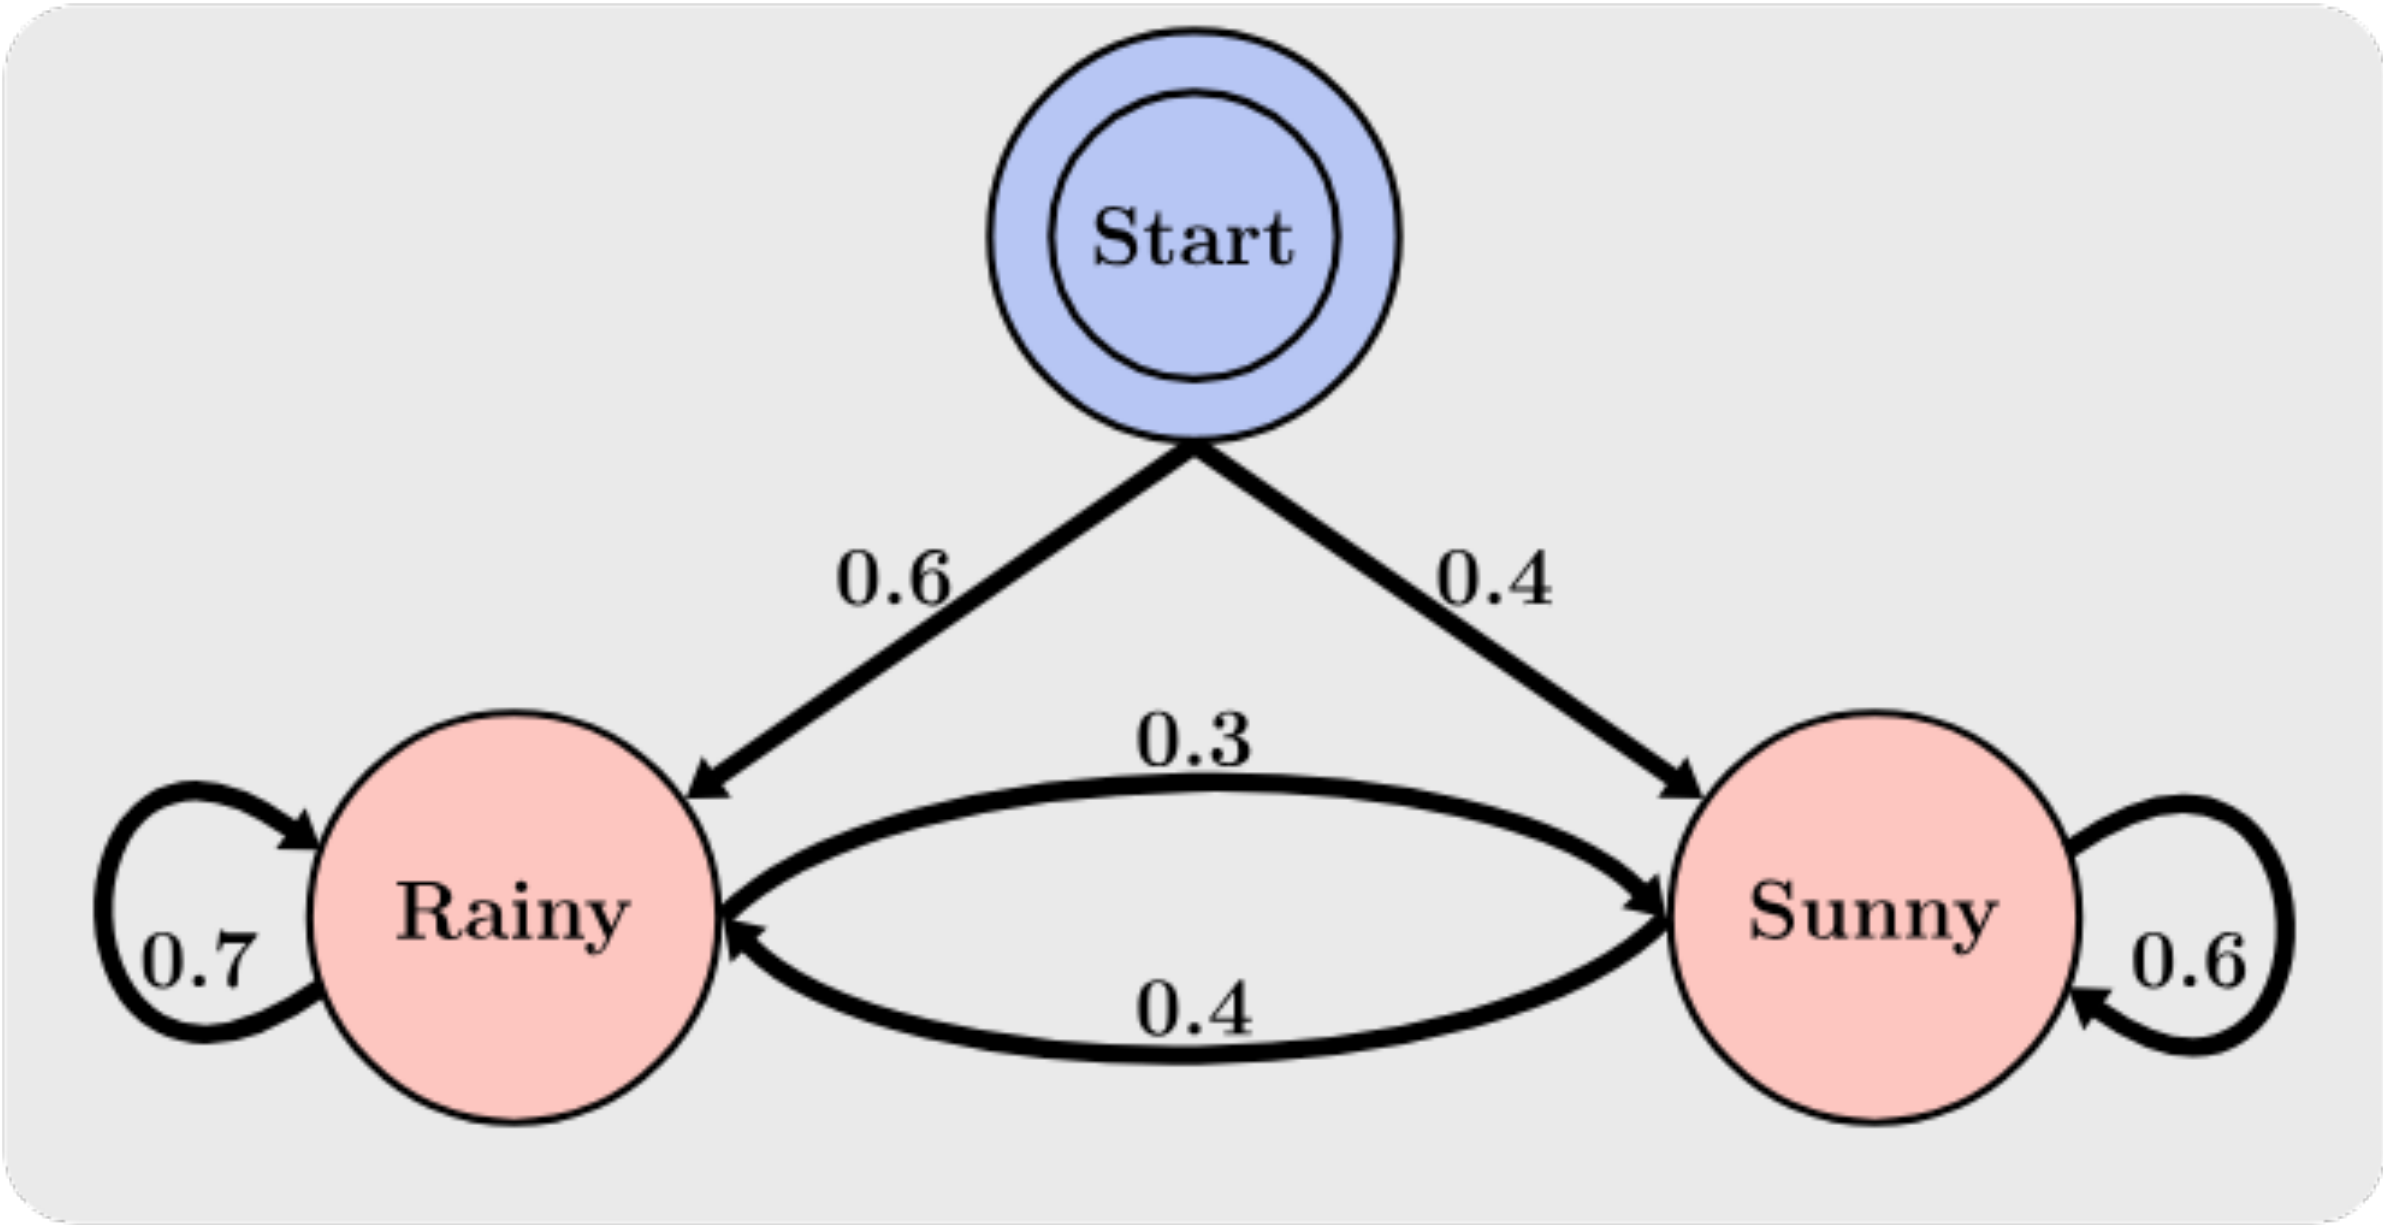
$$P(z_1, z_2, \dots, z_{t+1}) = P(z_1)P(z_2 | z_1) \dots P(z_t | z_{t-1})P(z_{t+1} | z_t)$$

Initial Conditions

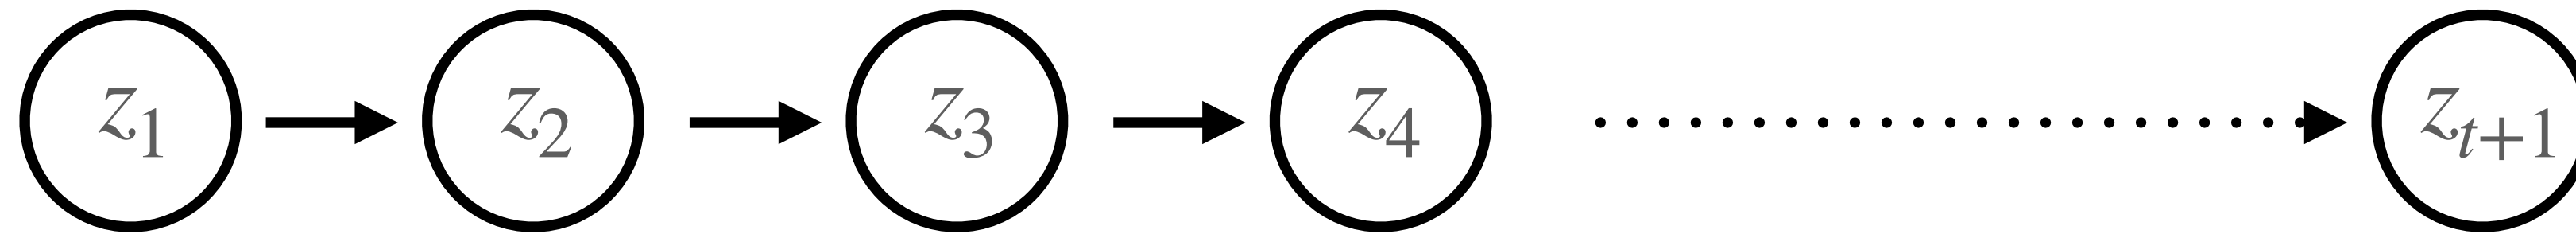
$$P(z_1 = R) = 0.6, \quad P(z_1 = S) = 0.4$$

Transitions Matrix

$P(z_{t+1} = R z_t = R) = 0.7$	$P(z_{t+1} = S z_t = R) = 0.3$
$P(z_{t+1} = R z_t = S) = 0.4$	$P(z_{t+1} = S z_t = S) = 0.6$

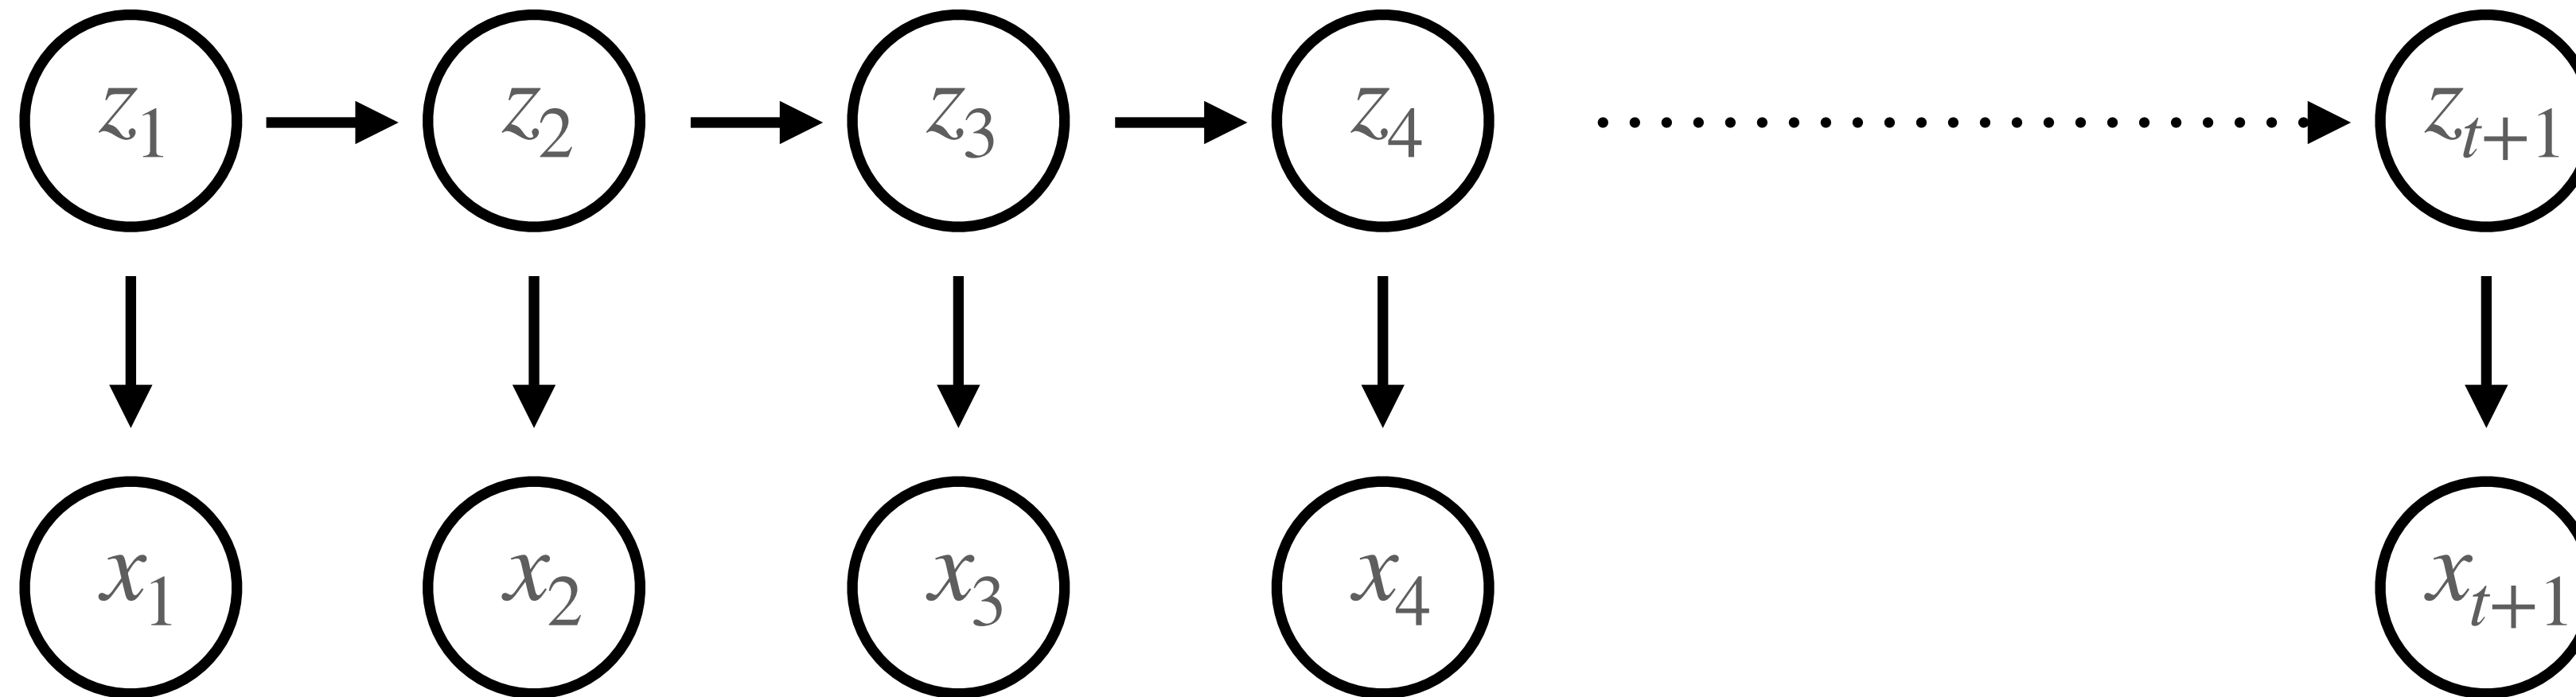


Hidden Markov Models




$$P(z_{1:T}) = P(z_1) \prod_{t=2}^T p(z_t | z_{t-1})$$

Hidden Markov Models



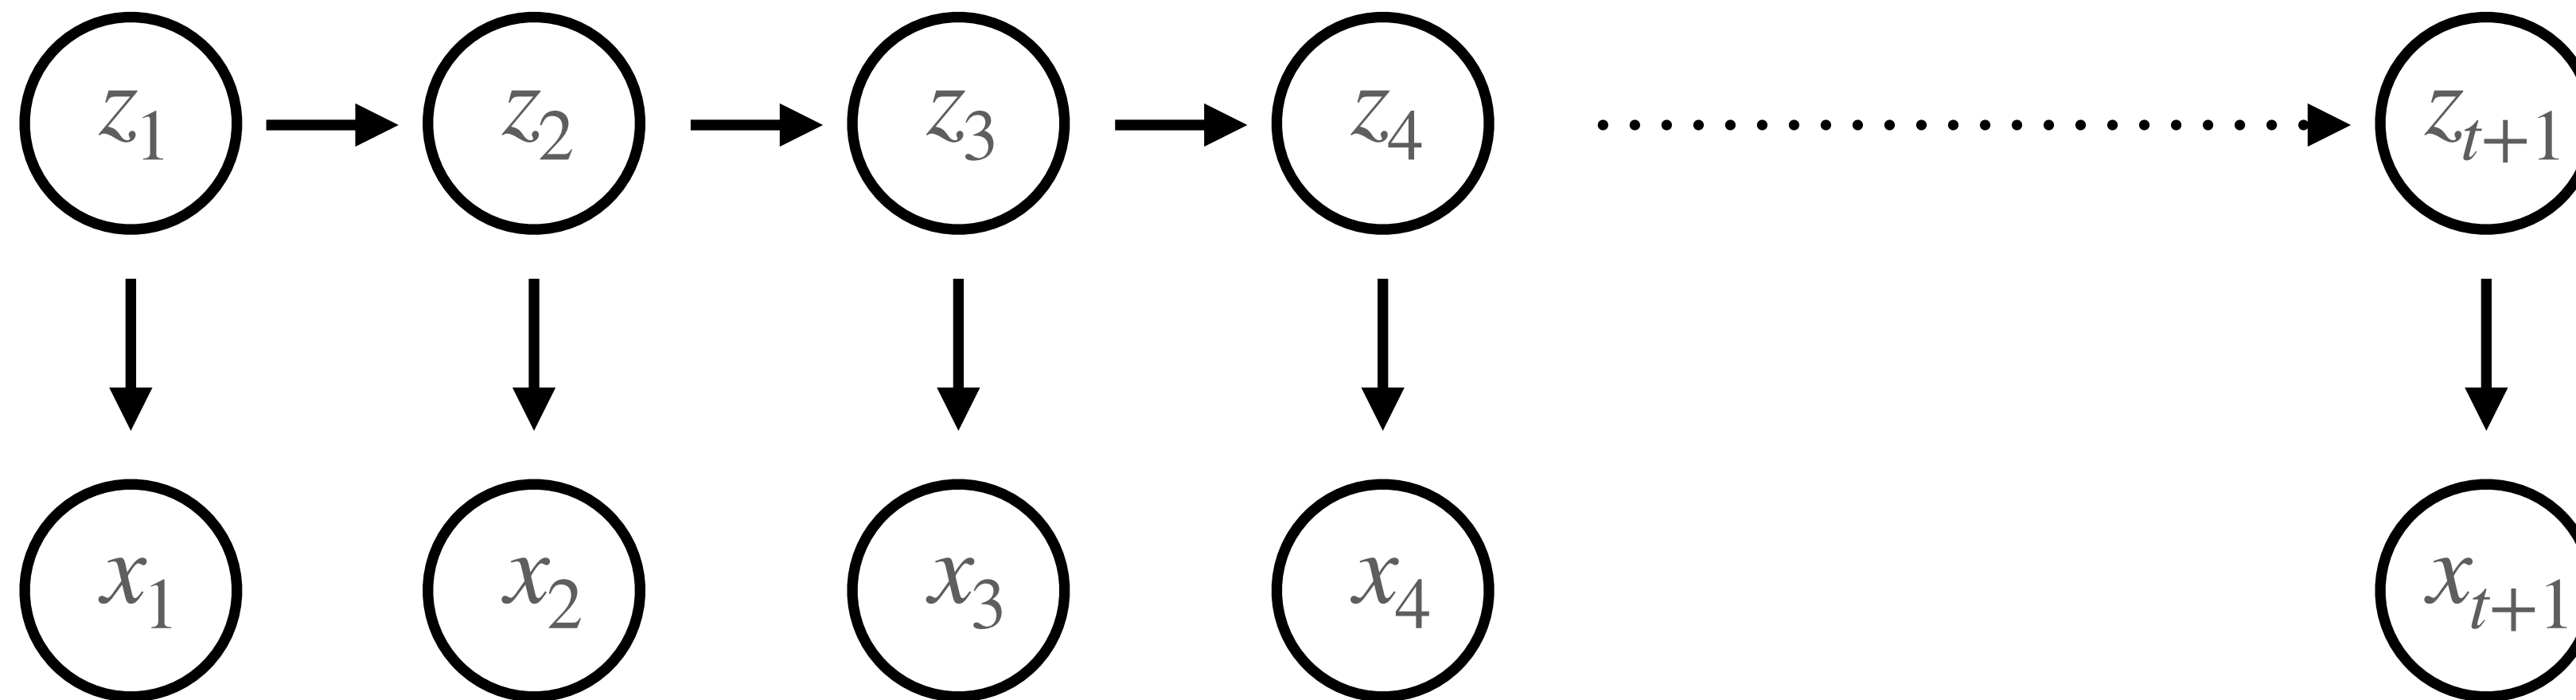
$$P(z_{1:T}) = P(z_1) \prod_{t=2}^T p(z_t | z_{t-1})$$

Hidden Markov Models



$$P(z_{1:T}, x_{1:T}) = P(z_1) \prod_{t=2}^T p(z_t | z_{t-1}) \prod_{t=1}^T p(x_t | z_t)$$

Hidden Markov Models



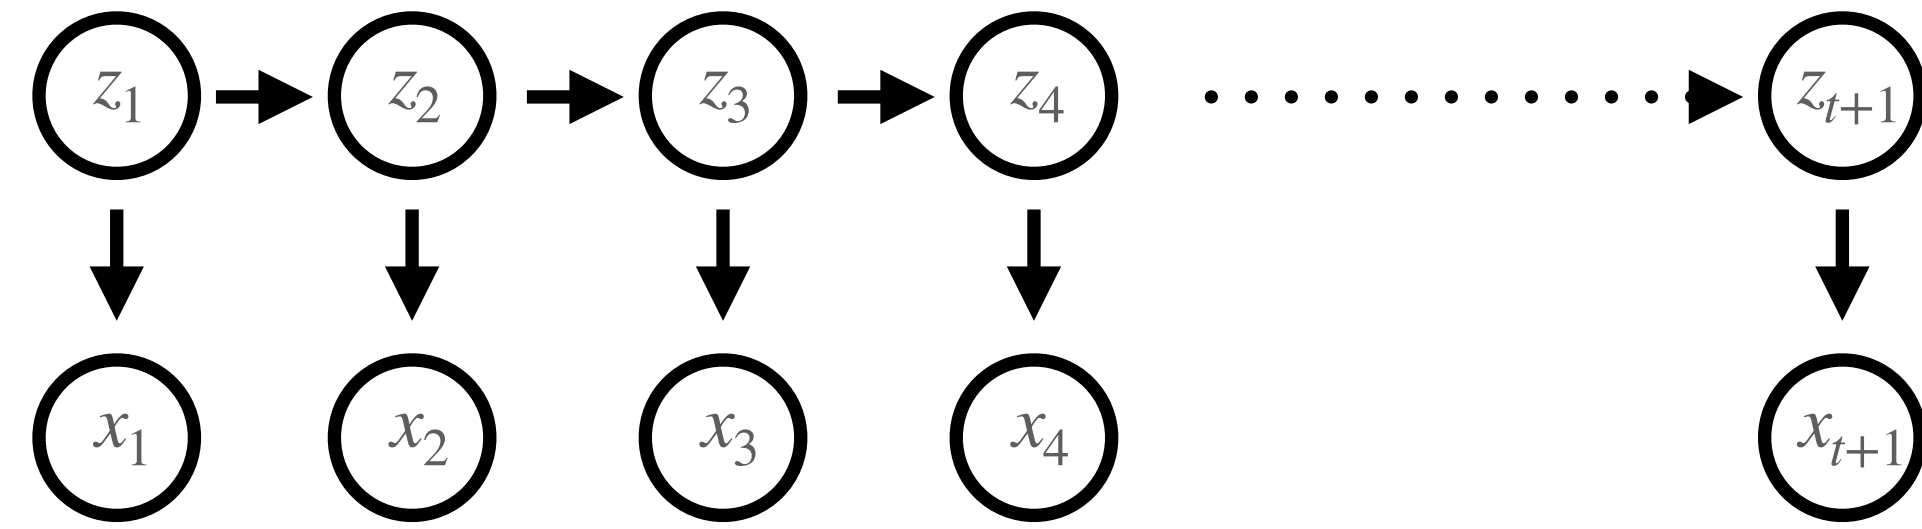
Initial
Probabilities

Transition
Probabilities

Emissions
(Observation)
Probabilities


$$P(z_{1:T}, x_{1:T}) = P(z_1) \prod_{t=2}^T p(z_t | z_{t-1}) \prod_{t=1}^T p(x_t | z_t)$$

Hidden Markov Models: Emissions Models




- $P(x|z)$ can take many different forms
 - Gaussian: $P(x_t|z_t) = \mathcal{N}(\mu_{z_t}, \sigma_{z_t})$

Hidden Markov Models: Emissions Models




- $P(x | z)$ can take many different forms
 - Gaussian: $P(x_t | z_t) = \mathcal{N}(\mu_{z_t}, \sigma_{z_t})$
 - Bernoulli, Poisson, etc.

Hidden Markov Models: Emissions Models



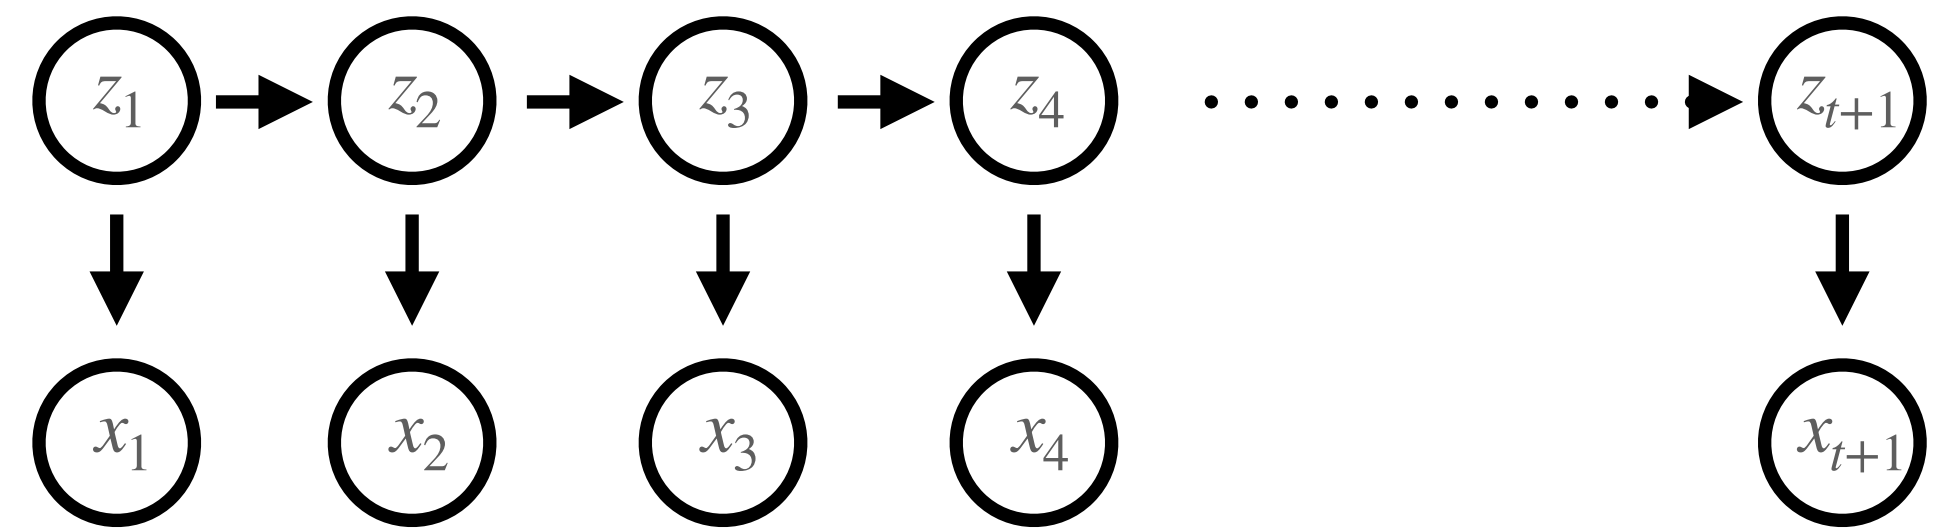
- $P(x|z)$ can take many different forms
 - Gaussian: $P(x_t|z_t) = \mathcal{N}(\mu_{z_t}, \sigma_{z_t})$
 - Bernoulli, Poisson, etc.
 - Autoregressive HMM (ARHMM):
 - Different dynamics in each discrete state: $P(x_t|z_t) = \mathcal{N}(y_{t-1} - A_{z_t}y_t, \sigma_{z_t})$

Hidden Markov Models: Emissions Models

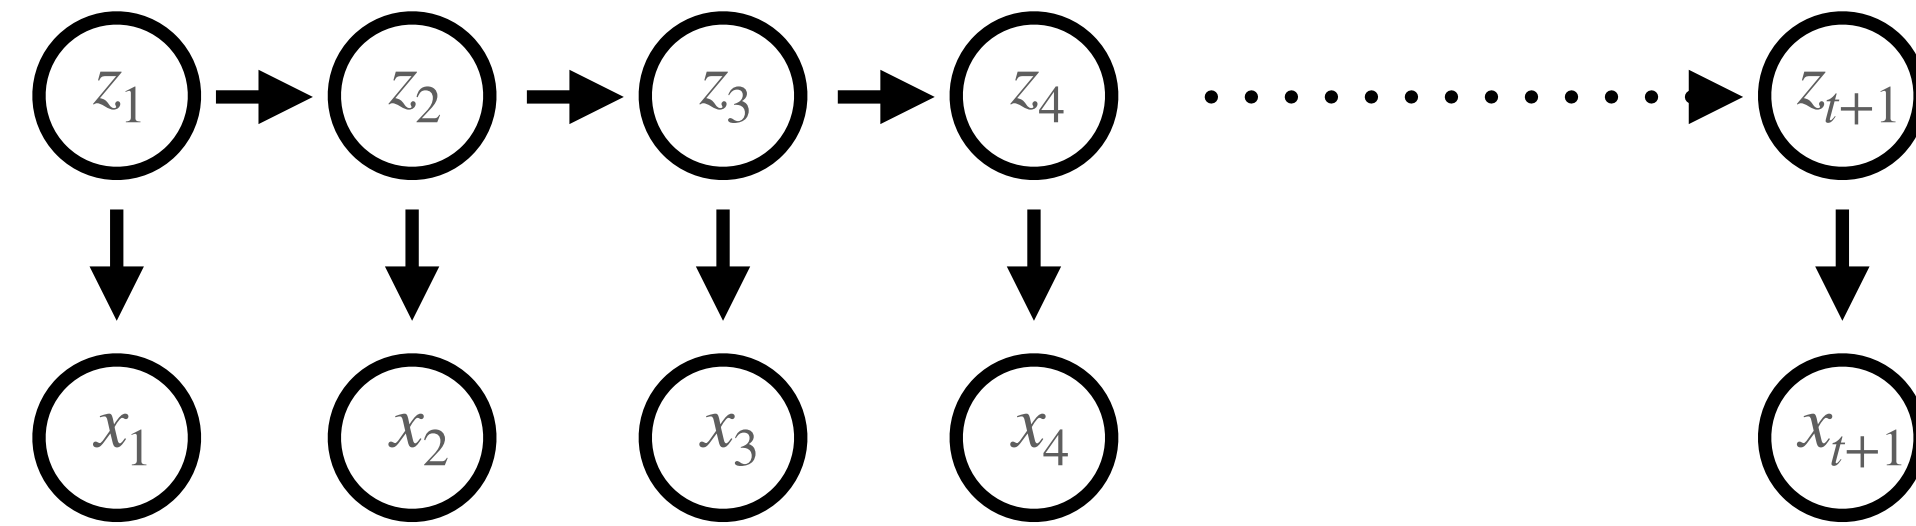


- $P(x|z)$ can take many different forms
 - Gaussian: $P(x_t|z_t) = \mathcal{N}(\mu_{z_t}, \sigma_{z_t})$
 - Bernoulli, Poisson, etc.
 - Autoregressive HMM (ARHMM):
 - Different dynamics in each discrete state: $P(x_t|z_t) = \mathcal{N}(y_{t-1} - A_{z_t}y_t, \sigma_{z_t})$
 - GLM-HMM:
 - Different GLM weights in each discrete state

Simulating from an HMM

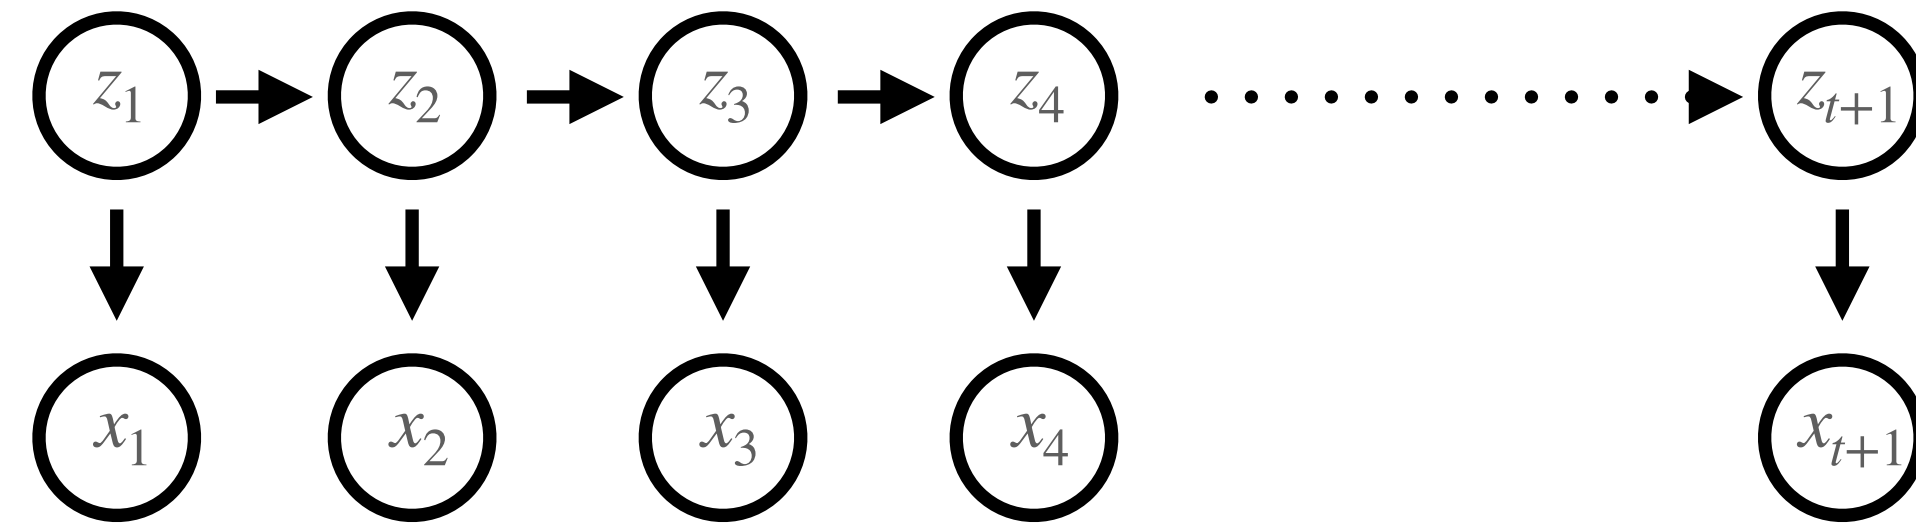


Simulating from an HMM



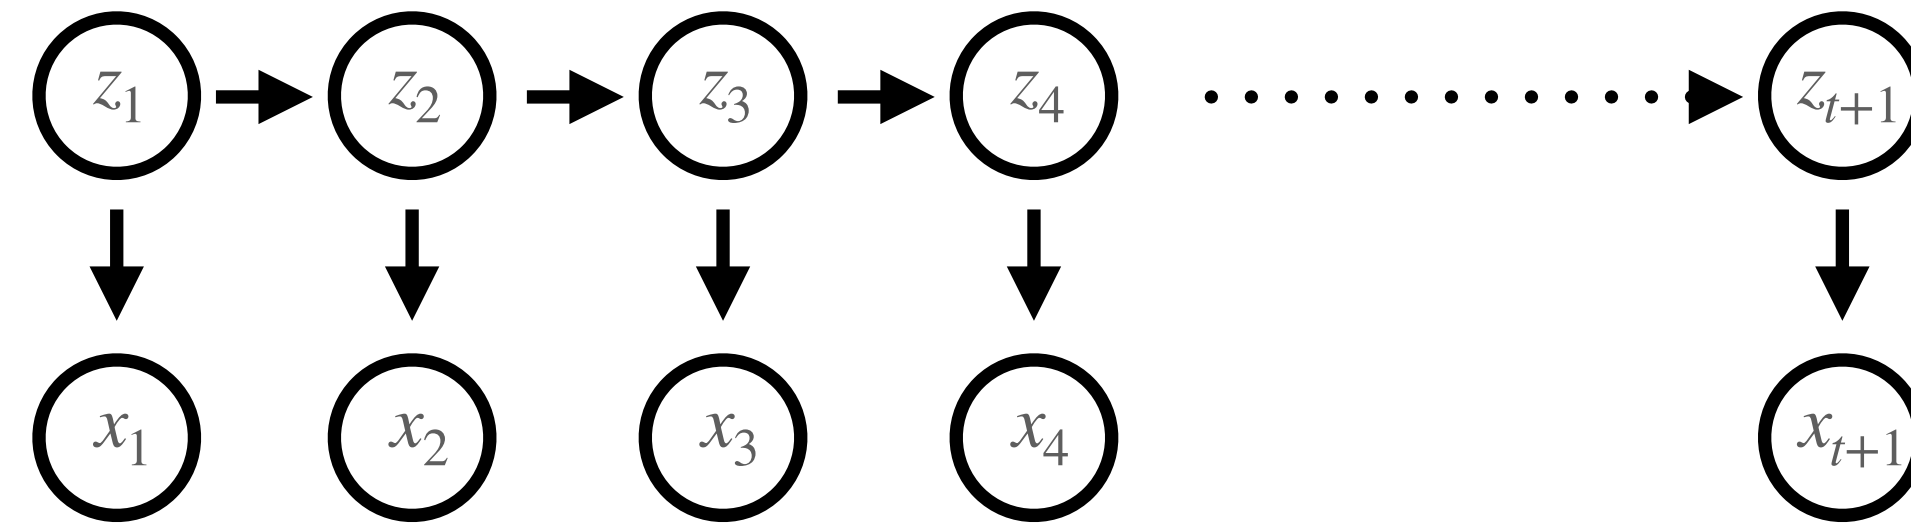
- Sample from $P(z_1)$

Simulating from an HMM



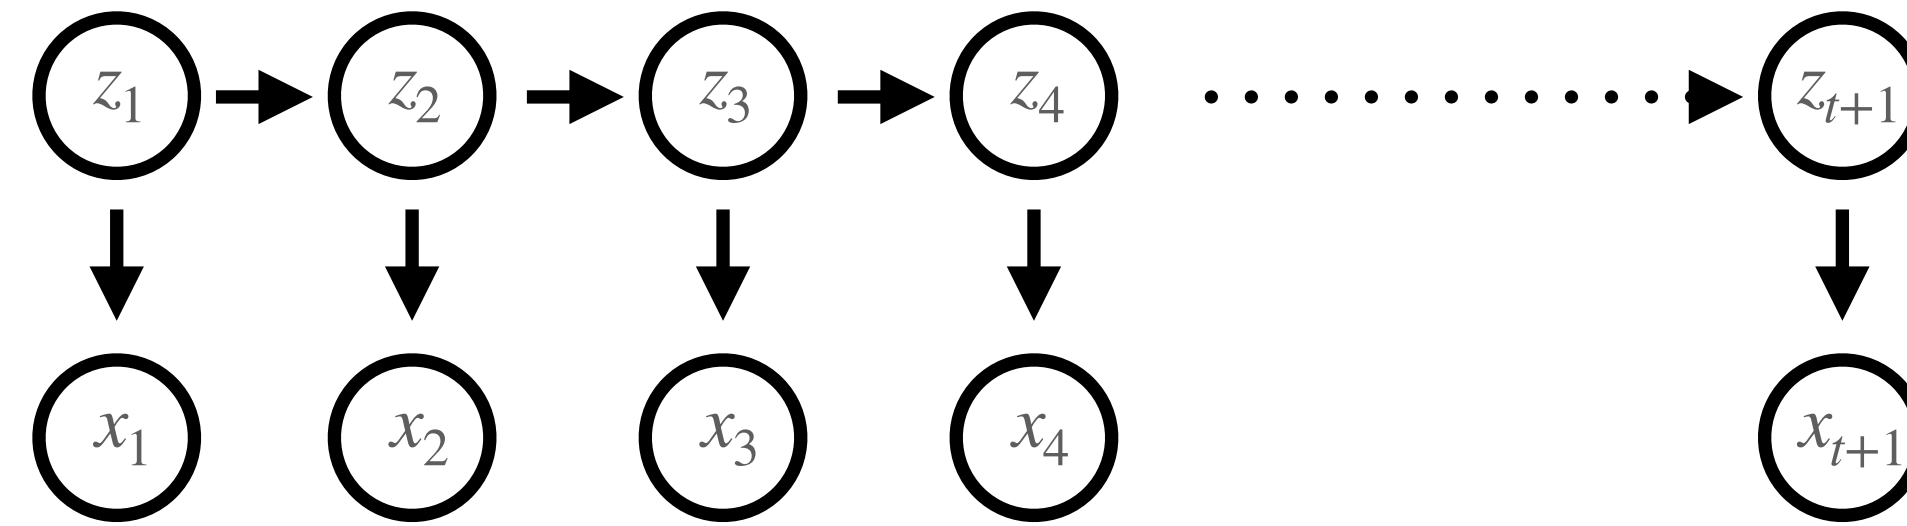
- Sample from $P(z_1)$
- Sample from $P(x_1 | z_1)$

Simulating from an HMM




- Sample from $P(z_1)$
- Sample from $P(x_1 | z_1)$
- For all future time steps:
 - Sample $P(z_{t+1} | z_t)$
 - Sample $P(x_{t+1} | z_{t+1})$

HMM: Goals



- Given some data:
 - Fit the model!
 - Infer discrete latent states with Forward/Backward Algorithm
 - Infer model parameters (transition probabilities, emissions model)
 - Determine the likelihood of the data given the model parameters

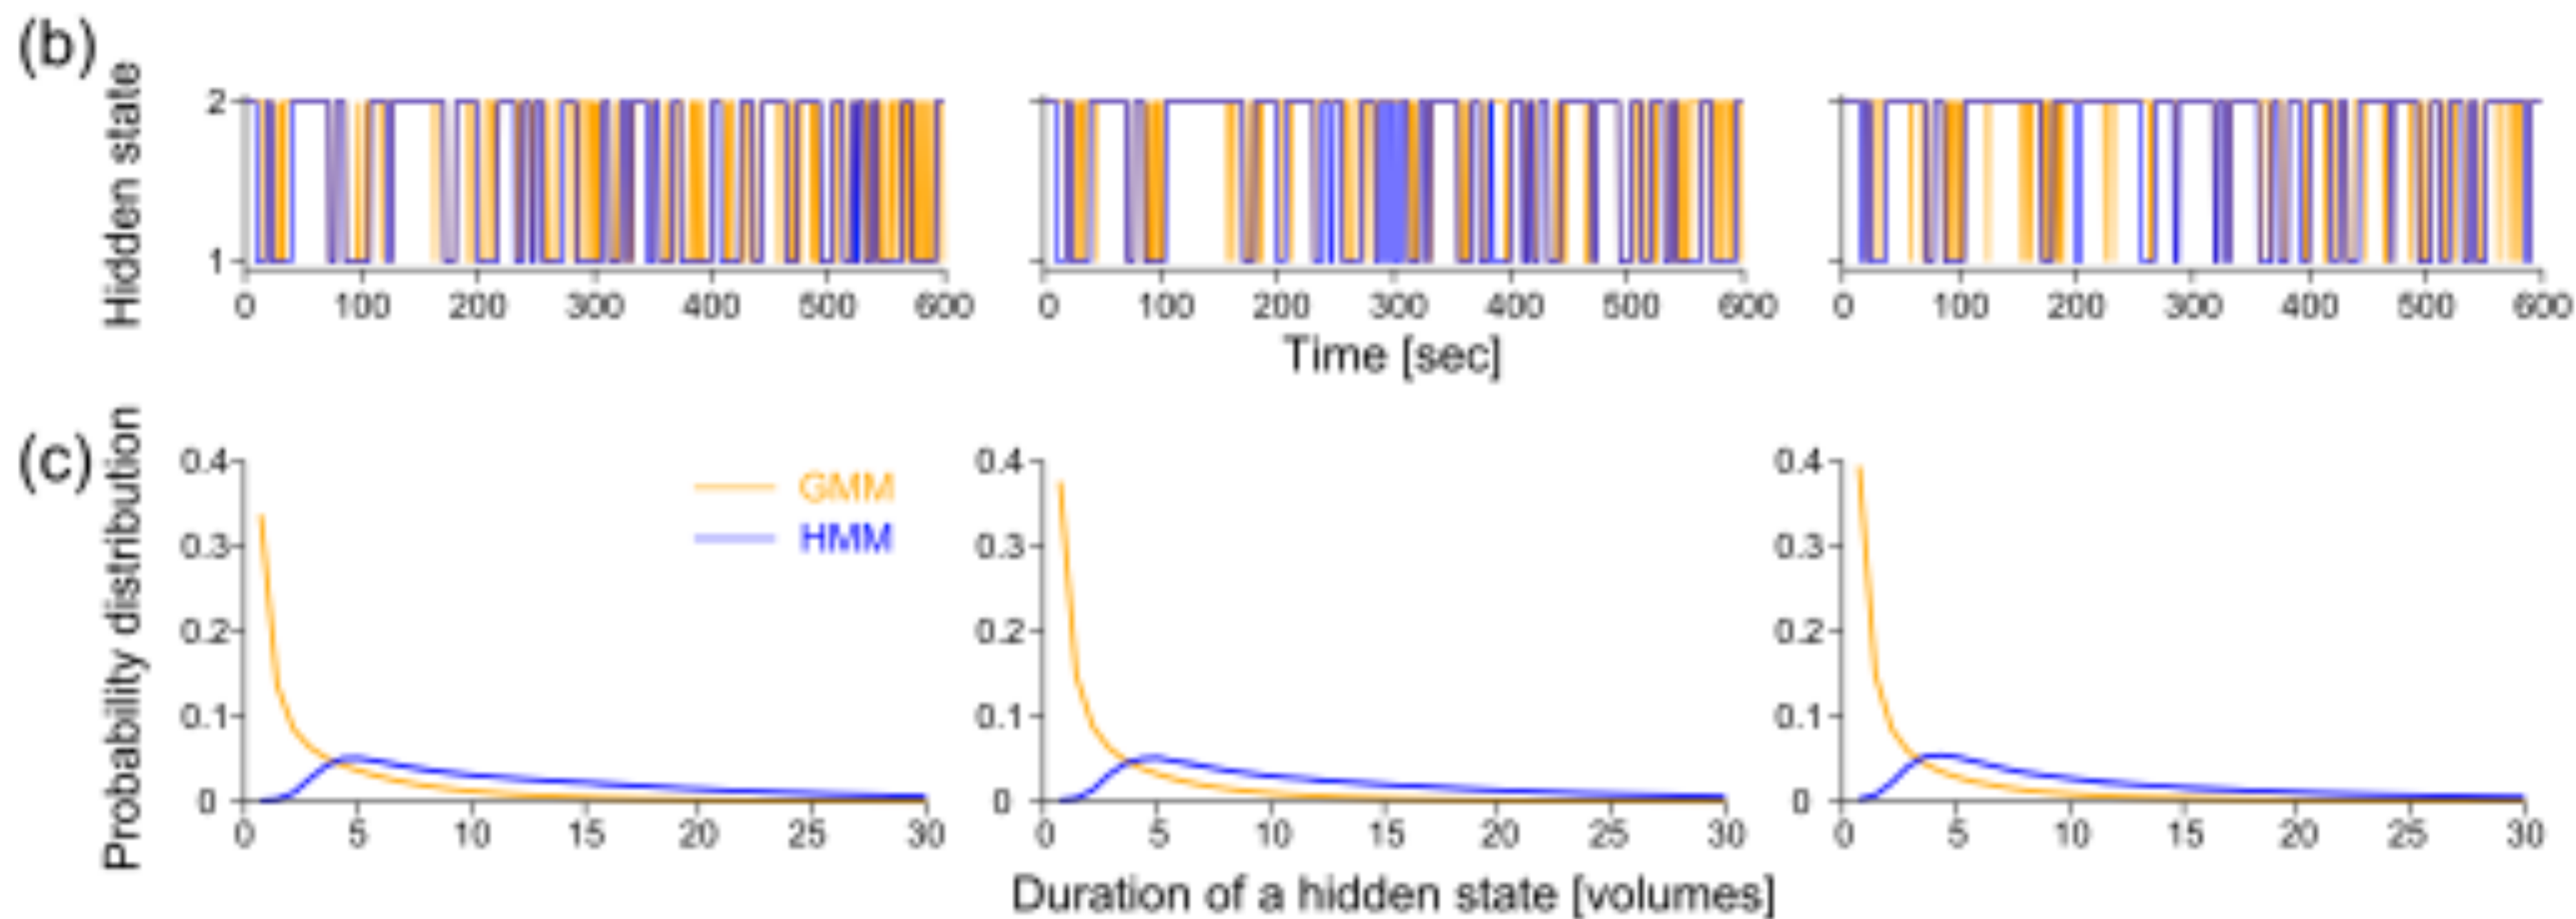
Hidden Markov Models: Emissions Models



- $P(x|z)$ can take many different forms
 - Gaussian: $P(x_t|z_t) = \mathcal{N}(\mu_{z_t}, \sigma_{z_t})$
 - Bernoulli, Poisson, etc.
 - Autoregressive HMM (ARHMM):
 - Different dynamics in each discrete state: $P(x_t|z_t) = \mathcal{N}(y_{t-1} - A_{z_t}y_t, \sigma_{z_t})$
 - GLM-HMM:
 - different GLM weights in each discrete state

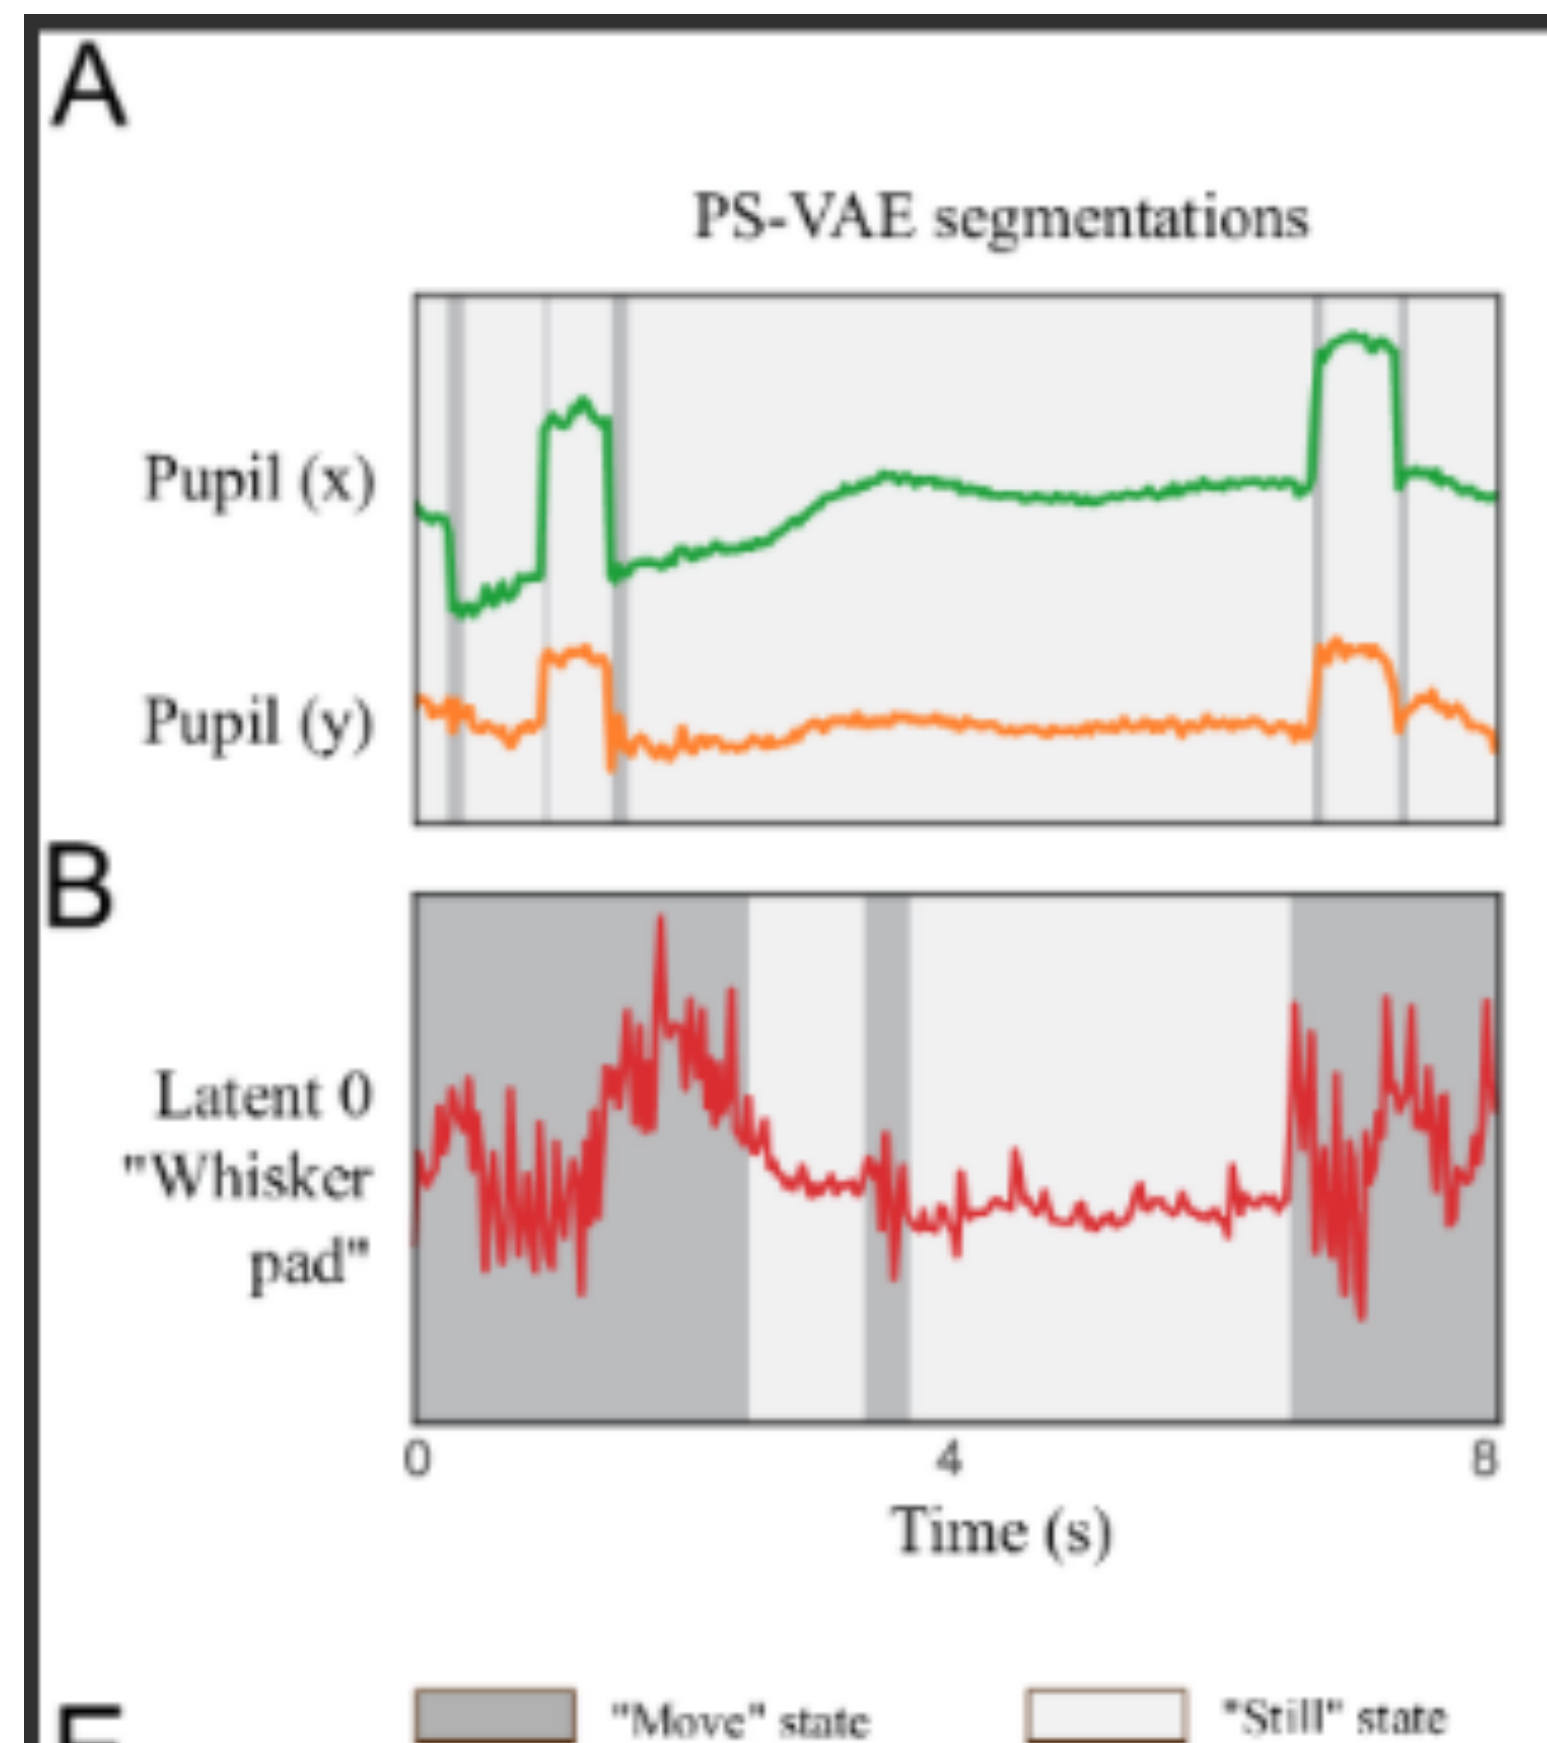
HMM - Gaussian Emissions

Modelling state-transition dynamics in resting-state brain signals by the hidden Markov and Gaussian mixture models



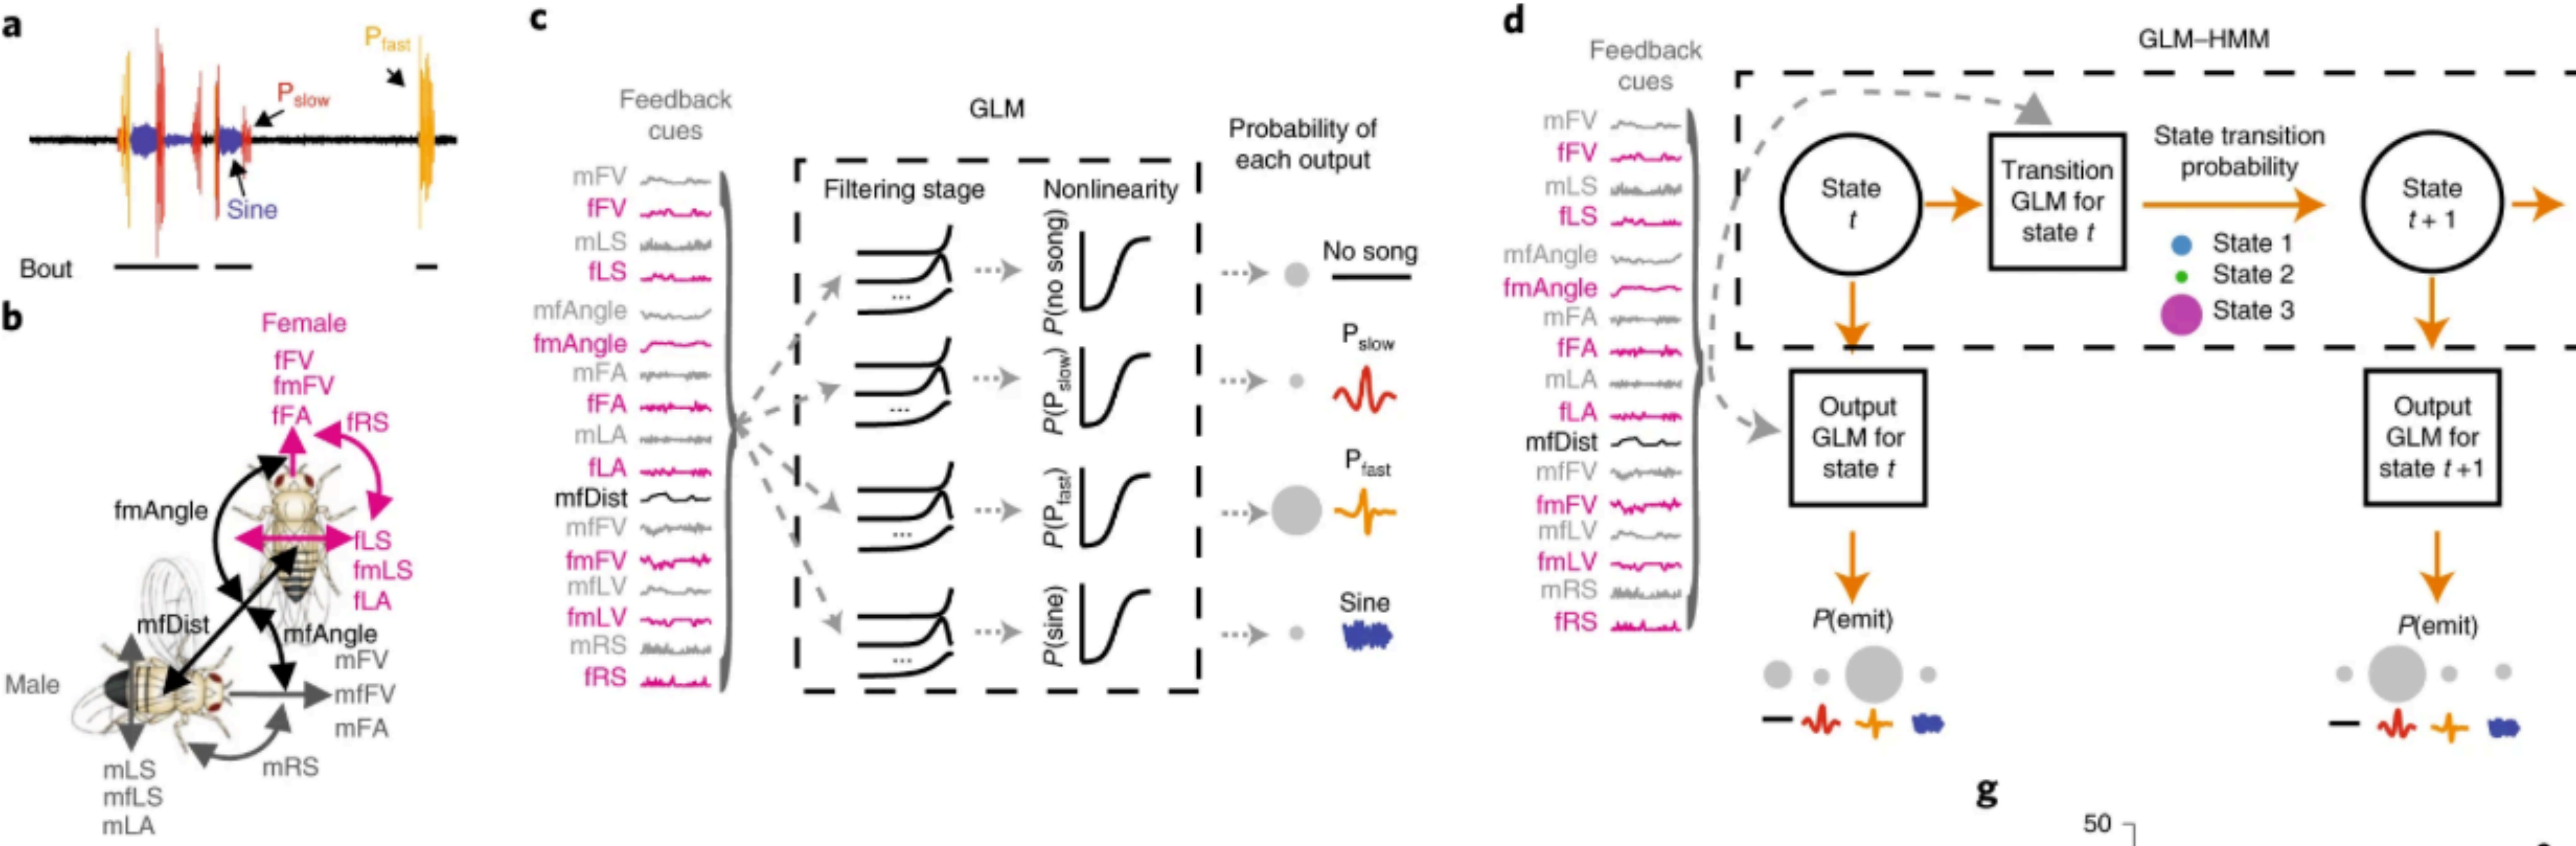
HMM - Autoregressive Emissions

Partitioning variability in animal behavioral videos using semi-supervised variational autoencoders





HMM - GLM Emissions

Unsupervised identification of the internal states that shape natural behavior



HMM - GLM Emissions

Mice alternate between discrete strategies during perceptual decision-making



Resources

- Probabilistic Machine Learning Book 1
 - <https://probml.github.io/pml-book/book1.html>
- Intro to LVM Notes from Princeton Course
 - https://pillowlab.princeton.edu/teaching/statneuro2020/notes/notes18_LatentVariableModels.pdf
- Interactive HMM Website
 - <https://nipunbatra.github.io/hmm/>